



CellML 1.0 Specification

This Version:

http://www.cellml.org/specifications/cellml_1_0

Latest Version:

<http://www.cellml.org/specifications>

Previous Versions:

<http://www.cellml.org/specifications/archive>

Authors

Warren Hedley (Bioengineering Research Group, University of Auckland)

Melanie Nelson (Physiome Sciences Inc.)

Contributors

David Bullivant (Bioengineering Research Group, University of Auckland)

Autumn Cuellar (Bioengineering Research Group, University of Auckland)

Yi Ge (Physiome Sciences Inc.)

Mark Grehlinger

Kam Jim (Physiome Sciences Inc.)

Scott Lett (Physiome Sciences Inc.)

David Nickerson (Bioengineering Research Group, University of Auckland)

Poul Nielsen (Bioengineering Research Group, University of Auckland)

Haoyu Yu (Physiome Sciences Inc.)

Abstract

This document specifies CellML™ 1.0, an XML-based language for describing and exchanging models of cellular and subcellular processes. MathML embedded in CellML documents is used to define the underlying mathematics of models. Models consist of a network of re-usable components, each with variables and equations manipulating those variables. Metadata may be embedded in CellML documents using RDF.

Status of this document

The 10 August 2001 version of the CellML 1.0 specification has been reviewed and endorsed by all of the authors of and contributors to the document. As of 10 August 2001, the syntax and semantics of all of the elements in the CellML 1.0 namespace are frozen.

The authors invite feedback from the public. Readers are encouraged to subscribe and send comments to the cellml-discussion mailing list. Alternatively, readers may send comments and questions via e-mail to info@cellml.org.

The latest version of the CellML specification is always available at the following URI:

- <http://www.cellml.org/specifications>

The list of errata associated with this document is available at the following URI:

- http://www.cellml.org/specifications/cellml_1_0/errata

Contents

- 1 Introduction
- 2 Fundamentals
- 3 Model Structure
- 4 Mathematics
- 5 Units
- 6 Grouping
- 7 Reactions
- 8 Metadata Framework
- A Using The CellML 1.0 DTD
- B Scripting functionality in CellML
- C Advanced Units Functionality
- D Changes

1 Introduction

1.1 Introduction to CellML

This document specifies CellML™ 1.0, an XML-based language for describing and exchanging models of cellular and subcellular processes. CellML is being developed by scientists at the University of Auckland (in the Bioengineering Research Group) and at Physiome Sciences, Inc. The development of CellML is guided by an advisory board drawn from many different areas of biological modelling (see the project team page on the CellML website for more information). CellML is being developed as an open standard, and all interested parties are encouraged to send feedback to info@cellml.org, or to the cellml-discussion mailing list.

1.1.1 Purpose and scope of CellML

CellML is intended to support the definition of models of cellular and subcellular processes. CellML facilitates the re-use of models and parts of models by using a component-based architecture. Models are split into logical

sub-parts called components that are connected together to form a model.

CellML separates the specification of the underlying mathematics of a model from a particular implementation of the model's solution. This makes a model independent of a particular operating system or programming language, and allows modellers to easily integrate parts of other peoples' models into their own models. CellML also allows the generation of equations for publishing from the same definition upon which the solution method is based, removing inconsistencies between the model and associated results in academic papers, and allowing others to reliably reproduce these results.

The scope of the CellML language is specifically limited to the definition of model structure. All other types of information that modellers need or want to include in a model document are incorporated using other languages. For instance, mathematics is included in CellML documents using Mathematical Markup Language (MathML). Metadata may be included using the Resource Description Framework and the Dublin Core Metadata Element Set.

1.1.2 What is XML?

The CellML language is defined in terms of a meta-language called eXtensible Markup Language (XML). XML is a standard published by the World Wide Web Consortium, the organisation responsible for defining many internet-related standards, most notably HTML. XML is essentially a means of adding structure to text documents, allowing machines to unambiguously associate text or binary data with a particular component in a document's data model.

XML is an appropriate medium for CellML because it is both human and machine readable. A model author can create a CellML document with a text editor or with CellML authoring software. XML is a well-defined and widely used specification. Many free software utilities and libraries for the processing of XML already exist, simplifying the development of CellML processing software. XML has also been designed to be usable over the internet, making CellML suitable for the interchange of models between software and databases at different locations.

A quick introduction to XML is available in the examples section of the CellML website.

1.1.3 Terminology

A model is a representation of the rules that govern the behaviour of a system. The terms in the following list provide two useful model classifications.

qualitative model

A model that defines the relationships between objects in the system, without defining any mathematics that represent the behaviour of those objects.

quantitative model

A model that defines the relationships between objects in the system, including the mathematics that represents the behaviour of those objects.

The terms defined in the following list are used in specifying the conformance of CellML documents and processing software to this specification.

may

Conforming CellML documents are permitted but not required to adhere to the limitation described.

Conforming CellML software is permitted but not required to behave as described.

must

Conforming CellML documents must adhere to the limitation described. Conforming CellML software must behave as described.

for interoperability

A non-binding recommendation included to increase the chances that CellML documents will be processed in a consistent manner by different applications.

error

A violation of the rules of this specification; results are undefined. Conforming CellML software may detect and report an error and may recover from it. The recommended best practice is for software to make information about errors available to the user.

valid CellML document

A document that conforms to all of the rules in this specification.

valid CellML subset document

A valid CellML document that only uses MathML elements from the CellML subset defined in Section 4.2.3.

CellML conformant software

CellML processing software that will interpret any valid CellML subset document according to the language semantics and processor rules defined in this specification.

fully MathML capable software

Software that can correctly interpret the full set of MathML content markup elements.

1.2 Structure of the CellML Specification

The CellML specification is divided into several sections, each of which discusses a particular aspect of CellML:

- Section 1 - **Introduction** - This section introduces CellML, XML, the terminology used throughout the specification, and the structure of the specification.
- Section 2 - **Fundamentals** - This section explains concepts used in all other sections of the specification, such as the definition of a valid CellML identifier and the use of XML namespaces in CellML.

- **Section 3 - Model Structure** - This section describes how models are organised in CellML. It includes an explanation of the use of a network of components to define a model and a discussion of variables in CellML.
- **Section 4 - Mathematics** - This section describes how mathematical expressions are defined in CellML documents using MathML, and defines the CellML subset of MathML elements.
- **Section 5 - Units** - This section explains the requirements for units in CellML and describes how a modeller can define arbitrary sets of units.
- **Section 6 - Grouping** - This section explains how a model can be organised into logical encapsulation and geometric containment hierarchies by grouping components.
- **Section 7 - Reactions** - This section introduces CellML syntax that allows the modeller to classify the involvement of the participants in the chemical expressions that make up reaction/pathway models.
- **Section 8 - Metadata Framework** - This section describes how RDF is used in CellML documents to define metadata and associate it with models, model components, and other CellML elements.
- **Appendices** - The appendices cover advanced and technical topics including the CellML DTD, recommendations for adding scripts to CellML documents, and units processing algorithms.

A valid CellML model can be created using nothing beyond the material covered in the fundamentals, model structure, mathematics, and units sections of the specification. The concepts in the remaining sections of the specification allow modellers to build more meaningful models.

Each section of the specification is further divided into five subsections:

- **Introduction** - This subsection explains the purpose of the elements covered in the current section.
- **Basic Structure** - This subsection describes the new elements and attributes introduced in the current section of the specification and how they are combined.
- **Examples** - This subsection provides one or two basic examples of the correct use of the elements and attributes introduced in the current section of the specification. More extensive examples can be found in the examples section of the CellML website.
- **Rules for CellML Documents** - This subsection provides formal rules for the use of the elements and attributes introduced in the current section to create valid CellML documents. These rules are specified as bulleted lists. Each rule may have an associated explanation, which appears directly after the rule, in square brackets ([]).
- **Rules for Processor Behaviour** - This subsection provides some rules for correct CellML processor behaviour with regards to the elements and attributes introduced in the current section.

Throughout the CellML specification, all XML elements and attributes that occur in the text are in the CellML namespace unless explicitly stated otherwise.

2 Fundamentals

2.1 Introduction

This section of the CellML specification introduces some concepts that are used throughout the entire language and defines rules that apply to all or many of the other parts of the specification. These include the definition of names and use of namespaces in CellML.

2.2 Basic Structure

2.2.1 Definition of a valid CellML identifier

The most common use of a CellML identifier is the `name` attribute required on many basic elements in CellML. The value of this attribute can be used to reference that element from elsewhere in the model definition or from another model definition altogether. An object's name can generally be thought of as a unique identifier for that object. Although the XML specification defines a mechanism for specifying that the value of an attribute is unique across an entire document (with the ID attribute type), this functionality is not used in CellML 1.0 because an object's name need only be unique across its own class of objects.

The generation of computer code for running simulations is one of the target applications for CellML. The value of an object's `name` attribute is intended to be a suitable name for the same object when it is represented in computer code. For this reason CellML identifiers must consist of only alphanumeric characters and the underscore character ("_"), and are subject to some additional constraints outlined below. These names will generally not be the most effective way of identifying the object to humans working with CellML models as it is not possible to include whitespace or formatting. More human readable names can be defined and associated with CellML objects using the metadata functionality introduced in Section 8.

The XML specification is based on the Unicode standard, which defines a scheme for 16 bit character encoding. Thus it is possible to include, for instance, Japanese characters in a valid XML document. In the interests of making the code generation process as convenient as possible for those using mainstream programming languages, CellML identifiers are subject to the following constraints:

- An identifier must consist only of alphanumeric characters from the US-ASCII character set and underscore characters,
- An identifier must contain at least one alphanumeric character.

Convenient code generation is also the reason why colons, periods, and hyphens may not appear in CellML identifiers. CellML identifiers are case sensitive: a variable with an identifier of `ABC` is different from a variable with an identifier of `abc`.

2.2.2 Namespaces in CellML

Namespaces in XML is a companion specification to the XML 1.0 specification. XML namespaces add a second level of naming to elements and attributes, allowing processing software to distinguish between elements and attributes from different languages. A namespace is identified by a Uniform Resource Identifier (URI), which has the feature of being unique. The value of a namespace URI need have nothing to do with the XML document that uses it. However, it typically points to a document that defines the rules for the language. The

URI may be mapped to a prefix, which may then be used in front of element and attribute names, separated by a colon. If not mapped to a prefix, the URI sets the default namespace for the current element and all of its children.

The CellML 1.0 specification defines a small number of elements and attributes and a namespace with which they must be associated. Putting CellML elements and attributes in the CellML namespace allows them to be distinguished from elements and attributes from other vocabularies with which CellML syntax might be combined in a CellML document. For instance, CellML makes use of the MathML vocabulary for the definition of equations, and all MathML elements must be placed in the MathML namespace in order for CellML processing software to recognise those elements. The use of namespaces also allows processing software to distinguish elements and attributes from different versions of the CellML specification. Applications that store their own proprietary data within a CellML document must define their own namespaces, and associate their own elements and attributes with those namespaces, as discussed in Section 2.2.3.

This specification is primarily concerned with the rules and semantics that relate to the elements and attributes in the CellML namespace, which are used in the definition of model structure. It is an error if documents contain elements and attributes in the CellML namespace that are not defined in this specification. This specification also defines how elements and attributes in the MathML, RDF and CellML Metadata namespaces can be combined with elements and attributes in the CellML namespace, and how processing software should deal with content in those namespaces. MathML is particularly important to CellML because content in this namespace is considered as fundamental as content in the CellML namespace. Metadata is defined using elements in the RDF namespace and linked to CellML elements using an `id` attribute in the CellML Metadata namespace as described in Section 8. Any CellML element may contain elements and attributes in other namespaces, which CellML processing software is free to ignore.

Table 1 lists the names, URIs and recommended prefixes of the namespaces referenced in this specification. For interoperability, the root element of any CellML document should set the default namespace and map the `cellml` prefix to the CellML 1.0 namespace URI. The latter simplifies the association of elements and attributes with the CellML namespace in regions of the document where the default namespace is not the CellML namespace. For instance, the MathML elements used to define equations are typically placed inside a `<math>` element that changes the default namespace to the MathML namespace. A `cellml:units` attribute in the CellML namespace can then be added to each of MathML's `<cn>` elements without having to redeclare the CellML namespace every time it is used.

Namespace Name	Namespace URI	Recommended Prefix
CellML	" http://www.cellml.org/cellml/1.0# "	<code>cellml</code>
CellML Metadata	" http://www.cellml.org/metadata/1.0# "	<code>cmeta</code>
MathML	" http://www.w3.org/1998/Math/MathML "	<code>mathml</code>
RDF	" http://www.w3.org/1999/02/22-rdf-syntax-ns# "	<code>rdf</code>

Table 1 The names, URIs and recommended prefixes of the namespaces referenced in this specification. See

text for more details.

2.2.3 Extending CellML documents

Any namespace with a URI not defined in Table 1 is an *extension namespace*. Any element in an extension namespace is an *extension element*. Any attribute in an extension namespace is an *extension attribute*. Model authors and CellML processing software may store information not covered by the CellML specification in a CellML document by defining their own extension elements and extension attributes. When authors and implementors define extension namespaces, it is recommended that they use URIs under their jurisdiction. Extension elements and extension attributes may appear anywhere in a CellML document as long as the result is well-formed XML.

For interoperability CellML processing software should respect the extension elements and attributes of other applications. If a model is created in application A, which adds its own extension elements, and is subsequently edited in application B, then application B should attempt to include application A's extension elements in its output, even if these extension elements are now invalid. Applications will need to validate their own extension data if a CellML document is read in from a non-trusted location.

The namespace extension mechanism provides a convenient way to associate a small amount of application-specific information with a model defined in CellML. However, it is recommended that applications needing to store large amounts of information, such as rendering or simulation information, do so in a separate document. This will make CellML documents easier to exchange, and will prevent the loss of application-specific information if the model is passed through applications unaware of the extensions.

2.3 Examples

Figure 1 contains some example CellML elements, each of which defines a `name` attribute. The values of the `name` attribute on the first three elements are valid CellML identifiers. The values of the `name` attribute on the last two elements are invalid identifiers.

```
<!--
The following elements have name attributes with valid values.
-->

<component name="my_favorite_component" />

<variable name="_ca2_conc" units="millimolar" />

<model name="model1345" />

<!--
```


The following elements have name attributes with invalid values.
Names may not consist purely of underscores or contain colons.
-->

```
<component name="___" />
```

```
<component name="my_model:my_component" />
```

Figure 1 XML elements defining name attributes. Valid and invalid CellML identifiers are shown, as noted in the comments.

Figure 2 contains portions of a typical CellML document that demonstrate the recommended use of namespaces. The root element sets the default namespace to the CellML namespace URI and explicitly maps the CellML namespace URI to the `cellml` prefix. The `<math>` element that encloses a set of equations inside a component element resets the default namespace to the MathML namespace. The `units` attribute on the `<cn>` element (which is in the MathML namespace) is placed in the CellML namespace by using the previously-defined `cellml` prefix.

```
<model
  name="simple_electrophysiological_model"
  xmlns="http://www.cellml.org/cellml/1.0#"
  xmlns:cellml="http://www.cellml.org/cellml/1.0#">

  ...

  <component name="extra_cellular_space">
    ...
    <math xmlns="http://www.w3.org/1998/Math/MathML">
      <apply><eq />
      <apply><diff />
      <bvar><ci> time </ci></bvar>
      <ci> Na </ci>
    </apply>
    <apply><times />
    <cn cellml:units="dimensionless"> -1.0 </cn>
    <ci> I_Na </ci>
  </apply>
</component>

  ...
```

```

</math>
</component>

...

</model>

```

Figure 2 A CellML fragment demonstrating the recommended use of namespaces in a CellML document. This fragment is taken from the simple electrophysiological model example on the CellML website.

Figure 3 demonstrates how software can embed its own information inside a valid CellML document using XML namespaces. The `<model>` element sets the default namespace to the CellML namespace, and maps the `app` prefix to an extension namespace (i.e., one not defined in Table 1). The `app` prefix is then used to define an `<app:component_rendering_information>` element and two attributes on a `<component>` element.

```

<model
  name="cellml_model_with_extensions"
  xmlns="http://www.cellml.org/cellml/1.0#"
  xmlns:app="http://www.physiome.org.nz/cellml_processor"
  xmlns:cellml="http://www.cellml.org/cellml/1.0#">

  <app:component_rendering_information>
    cell : blue
    membrane : yellow
    channel : red
  </app:component_rendering_information>

  <component
    name="cell"
    app:component_type="cell"
    app:render_corners="100, 100, 400, 400" />

</model>

```

Figure 3 A CellML document demonstrating the use of XML namespaces to embed application specific data inside a CellML document. The extension namespace URI was invented for demonstration purposes only.

2.4 Rules for CellML Documents

2.4.1 Valid CellML identifiers

- A valid CellML identifier must consist of only letters, digits and underscores, and must contain at least one letter or digit. This can be written using Extended Backus-Naur Form (EBNF) notation as follows:

```
letter    ::= 'a'...'z','A'...'Z'  
digit     ::= '0'...'9'  
identifier ::= ('_')* ( letter | digit ) ( letter | '_' | digit )*
```

[The variant of EBNF used above is defined in Section 6 of the XML 1.0 Recommendation.]

2.4.2 Proper use of the CellML namespace

- A document must not contain elements or attributes in the CellML namespace that are not defined in this specification. [Documents containing unknown elements or attributes in the CellML namespace are not valid CellML documents. Rules regarding the use of elements in the other namespaces defined in Table 1 are given in the appropriate sections. Note that attributes without an explicit prefix declaration are assumed to be in the same namespace as their parent element.]

2.4.3 Extension namespaces

- Although not explicitly stated throughout this specification, a document author may add extension elements and extension attributes to any CellML element in a CellML document without affecting the validity of the document. [Note that attributes without an explicit prefix declaration are assumed to be in the same namespace as their parent element.]
- For interoperability, elements in the CellML namespace should not be defined inside extension elements. [Specifically, applications should not define important model structure, mathematics or metadata information within extension elements, which other applications are free to ignore.]
- For interoperability, attributes in the CellML namespace should not be defined on extension elements.

2.4.4 Text nodes within CellML elements

- Any characters that occur immediately within elements in the CellML namespace must be either space (#x20) characters, carriage returns (#xA), line feeds (#xD), or tabs (#x9). [All of the elements in the CellML 1.0 namespace contain no text content. The characters listed above correspond to the definition of whitespace given in Section 2.3 of the XML Recommendation. Text content may still be included in extension elements inside CellML elements.]

2.5 Rules for Processor Behaviour

2.5.1 Treatment of CellML identifiers

- CellML processing software must handle identifiers in a case-sensitive manner. [Two CellML elements

of the same type may be defined with identifiers of **A** and **a**. Processing software is expected to match the identifiers in a case-sensitive manner when those elements are referenced at other places in the document.]

2.5.2 Treatment of attribute namespaces

- CellML processing software must treat attributes without an explicit namespace declaration as if they were in the same namespace as their parent element.

2.5.3 Treatment of extension namespaces

- CellML processing software may ignore extension elements and extension attributes. [If the namespace is unrecognised, then software should probably alert the user to its presence. Polite software should attempt to store non-CellML data so that it can write it out again when it exports the document. Software should validate its own non-CellML data carefully when reading documents from a non-trusted location.]
- CellML processing software may ignore the attributes and content of extension elements.

3 Model Structure

3.1 Introduction

Any model can be described as a network of connections between self-contained components. A component is a functional unit that may correspond to a physical compartment, event, or species or may be just a convenient modelling abstraction. A component contains variables and mathematical relationships that manipulate those variables. Connections exchange information between components. A connection contains mappings between variables in two components, allowing the value of a variable in one component to be passed to a variable in the other component.

3.2 Basic Structure

3.2.1 Definition of a model

A model is declared in CellML with a `<model>` element. This is the usual root element for a CellML document. The recommended best practice for specifying namespaces in a CellML document is described in Section 2.2.3.

The `<model>` element has a `name` attribute that allows this model to be unambiguously referenced by other models. For instance, this would be necessary if the model were to be combined with other models or partial models to create a larger model.

A `<model>` element may contain any number of the elements in the following list in any order. The recommended best practice is for elements placed within the `<model>` element to appear in the order given in the following list. This allows people to quickly find certain kinds of information within a CellML document.

- `<units>` - A modeller can declare a set of units to use in the model, as described in Section 5.
- `<component>` - Components are the smallest functional units in a model. Each component may contain variables that represent the key properties of the component and/or mathematics that describe the behaviour of the portion of the system represented by that component.
- `<group>` - Groups allow the modeller to define logical and physical relationships between components. Groups are defined using the `<group>` element, as discussed in Section 6.
- `<connection>` - Connections are used to connect components to each other and to map variables in one component to variables in another. Connections are defined using the `<connection>` element, as discussed in Section 3.2.4.

The `<model>` element (and indeed any of the elements in a CellML document) may define metadata to provide context for that object. This metadata might include documentation, citations from literature, or a modification history for the current CellML object. Adding metadata to a CellML document is discussed in detail in Section 8.

3.2.2 Definition of components

Constructing a model from multiple components encourages the re-use of components. For instance, an electrophysiological model of a cell might be organised into components that represent various ion channels. All of the mathematics that describe the behaviour of the L-type calcium channel would be defined in a single component representing this particular ion channel. If a modeller wished to re-use the portion of the model representing the L-type calcium channel in another model, he or she would only need to copy this component.

A `<component>` element is used to declare a CellML component. It must only be used inside a `<model>` element or as the root element of a CellML document. A `<component>` element that is the root of a CellML document does not define a complete model. It would probably be part of a library of standard components that could be imported and used in models. Eventually, CellML will include a mechanism that simplifies such re-use of components. At the present time, the component would need to be physically copied into a model document to be used in that model.

A CellML `<model>` may contain any number of `<component>` elements. Each `<component>` must have a `name` attribute, the value of which is a unique identifier for the component within the current model. The value of the `name` attribute is used to reference the component in other parts of the model, such as in connections and groups.

A `<component>` may contain any of the elements in the following list in any order. Again, recommended best practice is for elements placed within the `<component>` element to appear in the order given in the following list.

- `<units>` - A modeller can declare a set of units to use within the component, as described in Section 5.
- `<variable>` - A component may contain any number of `<variable>` elements, which define variables that may be mathematically related in the equation blocks contained in the component. Variables are

discussed in Section 3.2.3.

- `<reaction>` - A component may contain `<reaction>` elements, which are used to provide chemical and biochemical context for the equations describing a reaction. It is recommended that only one `<reaction>` element appear in any `<component>` element. The definition of reaction information is described in Section 7.
- `<mathml:math>` - A component may contain a set of mathematical relationships between the variables declared in this component. These equations are marked up using MathML, as discussed in Section 4. The `mathml` prefix is used to indicate that the `<math>` element is in the MathML namespace.

A `<component>` element is also a sensible place to define metadata, using the syntax presented in Section 8.

The definitions of two `<component>` elements are included in the example described in Section 3.3.

3.2.3 Definition of variables

Models are usually developed to investigate the behaviour of a number of variables that have biological significance. Each variable in the model belongs to a single component, which may contain equations that modify the value of that variable. The value of a variable may be passed through connections into other components. The variable must also be declared in these components, which can then use the value of the variable in their own equations but must not modify it.

The `<variable>` element is used to declare a CellML variable. It can only be used inside a `<component>` element. Variables must define a `name` attribute, the value of which must be unique across all variables in the current component. The name of a variable is used when referencing variables inside connections (see Section 3.2.4) and reactions (see Section 7). All variables must also define a `units` attribute. The value of this attribute must correspond to one of the keywords in the CellML units dictionary or the value of the `name` attribute of a `units` element defined within the current component or model, as described in Section 5.

A `<variable>` element may also have the following attributes:

- `initial_value` - This attribute provides a convenient means for specifying the value of a scalar real variable when all independent variables in the model have a value of 0.0. Independent variables are those with respect to which another variable is differentiated or integrated.
- `public_interface` - This attribute specifies the interface exposed to components in the parent and sibling sets (see below). The public interface must have a value "in", "out", or "none". The absence of a `public_interface` attribute implies a value of "none".
- `private_interface` - This attribute specifies the interface exposed to components in the encapsulated set (see below). The private interface must have a value "in", "out", or "none". The absence of a `private_interface` attribute implies a value of "none".

The name of the `initial_value` attribute derives from the fact that, in a model with only one independent variable, this would generally correspond to time, and so the value of the `initial_value` attribute sets the

starting condition for a simulation which progressed from time equals 0.0. The initial values of variables need not be set in the model definition at all. When multiple simulations are to be run using the same model, initial and boundary conditions are most conveniently set in an external simulation configuration file loaded separately by CellML processing software.

Whether or not a component may obtain the value of a variable in another component depends on the `public_interface` and `private_interface` attributes on the variable declaration, and the place of the two components in the encapsulation hierarchy. Encapsulation allows the modeller to hide a complex network of components from the rest of the model and provide a single component as an interface to the hidden network. Encapsulation effectively divides the network into layers, where connections between the layers must only be made through the interface components.

The components to which any given component may connect can be divided into four distinct sets with respect to any given component (the *current* component). The set of all components immediately encapsulated by the current component is referred to as the *encapsulated set*. If the current component is encapsulated, then the encapsulating component is referred to as the *parent*, and the set of all other components encapsulated by the same parent is referred to as the *sibling set*. If the current component is not encapsulated, then it has no parent and the sibling set consists of all other components in the model that are not encapsulated. All other components, which are not available to make connections with the current component, make up the *hidden set*. The encapsulation hierarchy and its effects on variable mapping are described in Section 6.

When a variable is declared with either a `public_interface` or `private_interface` attribute value of "in", then the value of that variable must be imported from another component. Otherwise, a variable's value must be set and modified in the current component. The variable is then said to *belong to* or be *owned by* the current component.

Eventually, it will be possible to specify the temporal and/or spatial variation of a variable's value using FieldML. The capability to include FieldML is still under development. At the present time, all variables must have scalar real values.

3.2.4 Definition of connections

Connections provide the mechanism for mapping variables declared within one component to variables in another component, allowing information to be exchanged between the various components in the network. The mapping of variables involves the transfer of a variable's value from one component to another, a process which may involve a conversion to ensure the units match. (More information on units conversion can be found in Section 5.)

The complete set of variable mappings between any two components constitutes a connection. Only one connection may be created between any given pair of components in a model. Each connection references the two components involved in the connection, and then maps variables from each of the components together. The interface attributes of each pair of variables must be compatible - an "out" variable in one component's interface must map to an "in" variable in the other component's interface. The direction of each mapping is

determined by the value of the `public_interface` and `private_interface` attributes on the two variables: the value is always passed from the variable with an interface value of "out" to the variable with an interface value of "in". The value of a variable declared with an interface value of "out" may be passed out to any number of variables in other components declared with interface values of "in". The component to which a variable belongs is found by tracing the variable back from "in" to "out" interfaces, following the model's connections.

The `<connection>` element is used to declare a CellML connection. It can only appear inside a `<model>` element.

A `<connection>` element must contain exactly one `<map_components>` element, which is used to reference the two components involved in the connection. Each `<map_components>` element must define `component_1` and `component_2` attributes, the values of which are the names of the components being referenced. In CellML 1.0, the referenced components must be defined within the current `<model>` element. It is anticipated that it will eventually be possible to reference components from other models, allowing models to be connected into larger models.

A `<connection>` element must also contain one or more `<map_variables>` elements, which are used to reference the variables being mapped between the two components in the connection. Each `<map_variables>` element must define `variable_1` and `variable_2` attributes, the values of which are equal to the names of variables defined in the components referenced by the `component_1` and `component_2` attributes on the `<map_components>` element, respectively. It is not necessary for the variables that are to be mapped to each other to have the same name, although this will typically be the case.

The CellML example discussed in Section 3.3 demonstrates the definition of a `<connection>` element.

3.3 Examples

Figure 4 contains a portion of the CellML description of the Hodgkin-Huxley squid axon model published in 1952. The excerpt contains the definitions of the components corresponding to the membrane and the sodium channel, and the connection between the two components. Most of the complexity from the full model definition has been left out for conciseness and clarity. This example is only used to demonstrate the standard use of the `<component>`, `<variable>`, and `<connection>` elements.

```
<model
  name="hodgkin_huxley_model_excerpt"
  xmlns="http://www.cellml.org/cellml/1.0#"
  xmlns:cellml="http://www.cellml.org/cellml/1.0#"
  xmlns:cmeta="http://www.cellml.org/metadata/1.0#">

  <!--
  Units definitions which could be referenced from the <variable> elements
```


would typically be inserted here. Units are discussed in Section 5.

-->

```
<component name="membrane">
```

```
<!-- the following variable is used in other components -->
```

```
<variable
```

```
name="V" initial_value="-75.0"
```

```
public_interface="out" units="millivolt" />
```

```
<!-- the following variables are imported from other components -->
```

```
<variable name="time" public_interface="in" units="millisecond" />
```

```
<variable name="i_Na" public_interface="in" units="microA_per_cm2" />
```

```
<variable name="i_K" public_interface="in" units="microA_per_cm2" />
```

```
<variable name="i_L" public_interface="in" units="microA_per_cm2" />
```

```
<!-- the following variable is only used internally -->
```

```
<variable name="C" initial_value="1.0" units="microF_per_cm2" />
```

```
<math xmlns="http://www.w3.org/1998/Math/MathML">
```

```
<apply id="membrane_voltage_diff_eq"><eq />
```

```
<apply><diff />
```

```
<bvar><ci> time </ci></bvar>
```

```
<ci> V </ci>
```

```
</apply>
```

```
<apply><divide />
```

```
<apply><minus />
```

```
<ci> i_Na </ci>
```

```
<ci> i_K </ci>
```

```
<ci> i_L </ci>
```

```
</apply>
```

```
<ci> C </ci>
```

```
</apply>
```

```
</apply>
```

```
</math>
```

```
</component>
```

```
<component name="sodium_channel">
```

```
<!-- the following variables are used in other components -->
```

```
<variable name="i_Na" public_interface="out" units="microA_per_cm2" />
```

```

<!-- the following variables are imported from other components -->
<variable name="time" public_interface="in" units="millisecond" />
<variable name="V" public_interface="in" units="millivolt" />

<math xmlns="http://www.w3.org/1998/Math/MathML">
<apply id="i_Na_calculation"><eq />
<ci> i_Na </ci>
... <!-- a function of V & time -->
</apply>
</math>
</component>

<connection>
<map_components component_1="membrane" component_2="sodium_channel" />
<map_variables variable_1="V" variable_2="V" />
<map_variables variable_1="i_Na" variable_2="i_Na" />
</connection>

</model>

```

Figure 4 A small portion of the CellML description of the Hodgkin-Huxley squid axon model from 1952. This excerpt contains the definition of the components corresponding to the membrane and the sodium channel, and the connection between them. Much detail has been omitted, but this example clearly demonstrates the relationship between the `<component>`, `<variable>` and `<connection>` elements.

The `membrane` component declares six variables, which are divided into three categories. The first variable is called `V`, and it represents the membrane voltage in the model. It has a `public_interface` attribute value of `"out"`, which indicates that the variable "belongs" to this component and that its value may be obtained by other components in the model via connections. It references a units definition by the name of `millivolt` (this definition is not included here), and is given an initial value of `-75.0` millivolts.

The subsequent four variables are `time`, `i_Na` (sodium current), `i_K` (potassium current) and `i_L` (leakage current). They are all declared with a `public_interface` attribute value of `"in"`, which indicates that their values are obtained from other components via connections.

Finally, a variable `C` (capacitance) is declared. This `<variable>` element defines neither a `public_interface` or a `private_interface` attribute. Both of these attributes therefore assume the default value of `"none"`, which means that the variable belongs to the current component and is not visible to other components in the model.

After the variable declarations, a `<math>` element in the MathML namespace is used to define an equation

relating V to the other variables. Only the values of the variables belonging to a component may be mathematically modified in that component. The equation included in Figure 4 is the well known differential equation from the Hodgkin Huxley model:

$$\frac{dV}{d\text{time}} = \frac{-(i_{Na} + i_K + i_L)}{C} \quad (1)$$

The `sodium_channel` component declares three variables, all of which represent quantities that were also declared in the membrane component. The `i_Na` variable declared in this component has a `public_interface` attribute value of "out", indicating that the sodium current belongs to this component. The value of the sodium current is calculated in this component, although the actual math has been omitted.

Finally, a `<connection>` element references the membrane and `sodium_channel` components using a `<map_components>` element, and maps the V and `i_Na` variables in each component together, using two `<map_variables>` elements. The value of the `variable_1` attribute on each `<map_variables>` element references the corresponding variable in the membrane component, since this is the component referenced by the `component_1` attribute on the `<map_components>` element. Similarly, the values of the `variable_2` attributes reference variables in the `sodium_channel` component.

3.4 Rules for CellML Documents

The following are the rules for using the `<model>`, `<component>`, `<variable>`, `<connection>`, `<map_components>`, and `<map_variables>` elements.

3.4.1 The `<model>` element

1. Allowed use of the `<model>` element

- A `<model>` element must contain only the following elements, which may appear in any order:
 - `<units>`, `<component>`, `<group>`, and `<connection>` elements in the CellML namespace,
 - `<RDF>` elements in the RDF namespace.

[The recommended best practice is to define the child elements in the CellML namespace in the order stated above.]

- Each `<model>` element must define a `name` attribute.

2. Allowed values of the `name` attribute

- The value of the `name` attribute must be a valid CellML identifier as discussed in Section 2.2.1.

3.4.2 The `<component>` element

1. Allowed use of the `<component>` element

- A `<component>` element must contain only the following elements, which may appear in any order:
 - `<units>`, `<variable>` and `<reaction>` elements in the CellML namespace,

- `<math>` elements in the MathML namespace,
- `<RDF>` elements in the RDF namespace.

[The recommended best practice is to define the child elements in the CellML and MathML namespaces in the order stated above. Note that a `<component>` element must not appear inside another `<component>` element. Such nesting could be intended to indicate a logical encapsulation relationship, a geometric containment relationship, or some other relationship between the two components. There is no reason to assume that the nesting hierarchy produced for one type of relationship would be consistent with the hierarchy produced for other types of relationship. Therefore, CellML defines these relationships using the `<group>` element, rather than nesting of `<component>` elements.]

- Each `<component>` element must define a `name` attribute.

2. Allowed values of the `name` attribute

- The value of the `name` attribute must be a valid CellML identifier as discussed in Section 2.2.1.
- The value of the `name` attribute must be unique across all `<component>` elements contained in the parent `<model>` element.

3.4.3 The `<variable>` element

1. Allowed use of the `<variable>` element

- A `<variable>` element must contain only the following elements:
 - `<RDF>` elements in the RDF namespace.
- Each `<variable>` element must define a `name` attribute and a `units` attribute. It may also define `public_interface`, `private_interface`, and `initial_value` attributes.

2. Allowed values of the `name` attribute

- The value of the `name` attribute must be a valid CellML identifier as discussed in Section 2.2.1.
- The value of the `name` attribute of a `<variable>` element must be unique across all `<variable>` elements contained in the same `<component>` element. [Two variables in the same component must not have the same name. However, two variables in different components may have the same name, and a variable may have the same name as its parent component.]

3. Allowed values of the `units` attribute

- The value of the `units` attribute must either be one of the keywords defined in the standard dictionary or the value of the `name` attribute on a `<units>` element defined in the current component or model. [The dictionary and the units element are described in Section 5.]

4. Allowed values of the `public_interface` attribute

- If present, the value of the `public_interface` attribute must be "in", "out", or "none".
- If not present, its value defaults to "none".

5. Allowed values of the `private_interface` attribute

- If present, the value of the `private_interface` attribute must be "in", "out", or "none".
- If not present, its value defaults to "none".

6. Proper use of the `public_interface` and `private_interface` attributes

- A `<variable>` element must not define both `public_interface` and `private_interface` attributes with values equal to "in". [A variable's value must only be obtained via one mapping.]

7. Allowed values of the `initial_value` attribute

- If present, the value of the `initial_value` attribute must be a real number.
- The absence of an `initial_value` attribute implies nothing. [The absence of this attribute would usually mean either that the variable does not need an initial value or that this value will be supplied in a parameter file or by the user at the time simulations using the model are run.]

8. Proper use of the `initial_value` attribute

- An `initial_value` attribute must not be defined on a `<variable>` element with a `public_interface` or `private_interface` attribute with a value of "in". [These variables receive their value from variables belonging to another component.]

3.4.4 The `<connection>` element

1. Allowed use of the `<connection>` element

- A `<connection>` element must contain only the following elements, which may appear in any order:
 - `<map_components>` and `<map_variables>` elements in the CellML namespace,
 - `<RDF>` elements in the RDF namespace.
- Each `<connection>` element must contain exactly one `<map_components>` element.
- Each `<connection>` element must contain at least one `<map_variables>` element. [It does not make sense to define a connection that does not map variables together. This rule prevents software from using empty connections to imply information not defined in this specification.]

3.4.5 The `<map_components>` element

1. Allowed use of the `<map_components>` element

- A `<map_components>` element must contain only the following elements:
 - `<RDF>` elements in the RDF namespace.
- Each `<map_components>` element must define a `component_1` attribute and a `component_2` attribute.

2. Allowed values of the `component_1` attribute

- The value of the `component_1` attribute must equal the value of the `name` attribute of a `<component>` element contained within the current `<model>` element.

3. Allowed values of the `component_2` attribute

- The value of the `component_2` attribute must equal the value of the `name` attribute of a `<component>` element contained within the current `<model>` element.

4. Proper use of the `component_1` and `component_2` attributes

- The `component_1` and `component_2` attributes on a single `<map_components>` element must not have the same value. [A connection must link two different components.]
- Each `<map_components>` element contained within `<connection>` elements that are contained

within a given `<model>` element must define a unique pair of `component_1` and `component_2` attribute values. [There can only be one connection between any two components in a network. This prevents setting up inconsistent, circular, or duplicate variable mappings between any two components in the network. However, it does not prevent a model author from creating inconsistent mathematical relationships between the variables.]

3.4.6 The `<map_variables>` element

1. Allowed use of the `<map_variables>` element

- A `<map_variables>` element must contain only the following elements:
 - `<RDF>` elements in the RDF namespace.
- Each `<map_variables>` element must define a `variable_1` attribute and a `variable_2` attribute.

2. Allowed values of the `variable_1` attribute

- The value of the `variable_1` attribute must equal the value of the `name` attribute of a `<variable>` element contained in the `<component>` element referenced by the `component_1` attribute on the `<map_components>` element within the current `<connection>` element.

3. Allowed values of the `variable_2` attribute

- The value of the `variable_2` attribute must equal the value of the `name` attribute of a `<variable>` element contained in the `<component>` element referenced by the `component_2` attribute on the `<map_components>` element within the current `<connection>` element.

4. Proper use of the `<map_variables>` element to map variables to each other

[The rules for mapping a variable to other variables depend on the encapsulation hierarchy of the component that owns the variable. This hierarchy divides the rest of the components in the model into *parent*, *sibling*, *encapsulated*, and *hidden* sets, as described in Section 3.2.3. The `public_interface` attribute defines the availability of a variable to the parent component and components in the sibling set. The `private_interface` attribute defines the availability of a variable to components in the encapsulated set. Variables are not available to components in the hidden set.]

- Variables with a `public_interface` or `private_interface` attribute value of "in" must be mapped to variables with a `public_interface` or `private_interface` attribute value of "out".
- A variable with either a `private_interface` or `public_interface` attribute value of "in" must be mapped to no more than one other variable in the model. [Note that a similar restriction does not apply to variables with interface values of "out". An output variable can be mapped to multiple input variables in various components in the current model.]
- A variable with a `public_interface` attribute value of "in" must be mapped to a single variable owned by a component in the sibling set, provided the target variable has a `public_interface` attribute value of "out", or to a single variable owned by the parent component, provided the target variable has a `private_interface` attribute value of "out".
- A variable with a `public_interface` attribute value of "out" may be mapped to variables

owned by components in the sibling set, provided the target variables have `public_interface` attribute values of "in". It may also be mapped to variables owned by the parent component, provided the target variables have `private_interface` attribute values of "in".

- A variable with a `private_interface` attribute value of "in" may be mapped to a single variable owned by a component in the encapsulated set, provided the target variable has a `public_interface` attribute value of "out".
- A variable with a `private_interface` attribute value of "out" may be mapped to variables owned by components in the encapsulated set, provided the target variables have `public_interface` attribute values of "in".

3.5 Rules for Processor Behaviour

3.5.1 Mapping of variables

For interoperability, CellML processing software should take into account the units definitions referenced by any two variables that are mapped together. If the units references are not equivalent, as defined in Appendix C.2.1, then a conversion may be required. An algorithm for performing this conversion is proposed in Appendix C.3.5.

4 Mathematics

4.1 Introduction

CellML allows modellers to unambiguously specify the underlying mathematics of a cellular model. Model components may contain mathematical expressions that manipulate the values of variables that belong to them. These expressions are also free to use (but must not modify) the values of any other variable declared in those components.

Mathematical expressions are embedded in CellML documents using Mathematical Markup Language 2.0 (MathML), an XML-based language that encodes the underlying structure of a mathematical expression. CellML uses a subset of the elements from MathML 2.0, known as the *content markup* element set, which includes several deprecated elements from MathML 1.0.

CellML 1.0 does not require processing software to implement support for scripting. If software chooses to do so, some recommendations on the use of scripting are given in Appendix B.

4.2 Basic Structure

4.2.1 Definition of mathematics

All mathematical expressions defined using MathML must be placed inside a `<mathml:math>` element. `<mathml:math>` elements must only be defined in `<cellml:component>` or `<cellml:role>` elements. The `mathml` and `cellml` namespace prefixes are used throughout this section to indicate that elements are in

the MathML and CellML namespaces, respectively. The `<cellml:role>`, `<cellml:variable_ref>` and `<cellml:reaction>` elements mentioned in this section are described in detail in Section 7 of this specification.

`<mathml:math>` elements that occur as child elements of `<cellml:component>` elements can be used to define arbitrary expressions relating the variables declared in that component. A mathematical expression may make use of any variable declared within the current component by placing the variable's name within a `<mathml:ci>` element. Expressions must only modify the values of variables that *belong* to that component. Variables that belong to a component are those that are not declared with a `public_interface` or `private_interface` attribute value of "in".

`<mathml:math>` elements that occur as child elements of `<cellml:role>` elements (these are defined within `<cellml:variable_ref>` elements, which are in turn defined within `<cellml:reaction>` elements) can be used to define expressions that modify the values of specific variables in specific ways. These expressions may make use of any variable declared in the current component but must only modify the value of the variable referenced by the ancestor `<cellml:variable_ref>` element, subject to further limitations that are described in Section 7.2.

CellML processing software must interpret MathML elements according to the semantics defined in the MathML 2.0 Recommendation. However, CellML 1.0 does define some restrictions on, and additions to, the MathML syntax. These are covered in the subsequent sections.

4.2.2 MathML's presentation and content markup elements

The complete set of elements defined in the MathML 2.0 Recommendation is split into two principal sub-vocabularies: the *presentation markup* and *content markup* elements. The presentation markup elements describe the visual rendering of mathematical expressions and objects. The content markup elements specify the underlying meaning of a mathematical expression or object, without regard to its presentation.

CellML is used to describe the structure and mathematics of cellular models. For this reason, valid CellML documents must only contain content markup elements within a `<mathml:math>` element. There is one exception: model authors may associate rendering information with a particular expression by placing MathML presentation markup elements inside a `<mathml:annotation-xml>` element. CellML processing software may ignore the contents of `<mathml:annotation>` and `<mathml:annotation-xml>` elements.

An example demonstrating the embedding of MathML content and presentation markup elements in a CellML document is presented in Section 4.3.

4.2.3 The CellML subset of MathML content elements

Valid CellML documents may contain any MathML content markup elements within a `<mathml:math>` element, as long as the arrangement of these elements follows the rules defined in the MathML 2.0 Recommendation. However, it is anticipated that it will be some time before software is able to interpret all of these elements. To encourage interoperability, this section defines a subset of the MathML content markup

elements known as the *CellML subset*. CellML documents that only contain content markup elements from the CellML subset are known as *valid CellML subset documents*. CellML processing software may only call itself *CellML conformant* if it is able to correctly interpret all of the MathML elements in the CellML subset according to the semantics defined in the MathML 2.0 Recommendation.

The complete list of MathML elements in the CellML subset is given in Figure 5. Many of the elements in the CellML subset are included to provide facilities for the definition of algebraic and ordinary differential equations. Others (such as the trigonometric operators) have been included because they are reasonably straightforward to translate to computer code.

token elements

<cn>, <ci>

basic content elements

<apply>, <piecewise>, <piece>, <otherwise>

relational operators

<eq>, <neq>, <gt>, <lt>, <geq>, <leq>

arithmetic operators

<plus>, <minus>, <times>, <divide>, <power>, <root>, <abs>, <exp>, <ln>, <log>, <floor>, <ceiling>, <factorial>

logical operators

<and>, <or>, <xor>, <not>

calculus elements

<diff>

qualifier elements

<degree>, <bvar>, <logbase>

trigonometric operators

<sin>, <cos>, <tan>, <sec>, <csc>, <cot>, <sinh>, <cosh>, <tanh>, <sech>, <csch>, <coth>, <arcsin>, <arccos>, <arctan>, <arccosh>, <arccot>, <arccoth>, <arccsc>, <arccsch>, <arcsec>, <arcsech>, <arcsinh>, <arctanh>

constants

<true>, <false>, <notanumber>, <pi>, <infinity>, <exponentiale>

semantics and annotation elements

<semantics>, <annotation>, <annotation-xml>

Figure 5 The *CellML* subset of MathML content markup elements, grouped according to function. All elements in this figure are in the MathML namespace.

4.2.4 Ordering of expressions

The mathematics in a model defined using CellML 1.0 consist of a static system of expressions, which are distributed over a network of components. CellML does not define the order of evaluation of equations, as this is simulation information rather than model information.

4.2.5 Scope of expressions

Within a CellML model, all expressions are assumed to have unlimited scope with respect to the independent variables unless explicitly stated using MathML's `<piecewise>` construct or some other form of conditional expression. This means that if the initial conditions for a variable, the value of which is determined by a differential equation, are to be specified using an equality, the two equations should have their scope limited so that they do not contradict each other.

4.2.6 Associating units with numbers

To ensure that models are robust and portable, all variables and numbers that occur in mathematical expressions within a CellML document must have units associated with them. CellML's units framework is introduced in Section 5 and the association of units with variables is presented in Section 3.2.3. The association of units with numbers in equations requires an extension to MathML. This can be done in a manner consistent with the association of units with variables and with application-specific extensions to CellML by adding a `units` attribute in the CellML namespace to the `<mathml:cn>` element, which encloses all numbers. The example presented in Section 4.3 demonstrates this.

4.2.7 Definition of scripts

CellML 1.0 does not define a standard method by which model authors can embed scripts in CellML documents in a portable way. It is anticipated that this functionality will be defined in a subsequent version of CellML. However, the use of scripts in CellML is strongly discouraged. CellML is aimed at specifying a model in terms of its most basic governing equations. Wherever possible, mathematical equations should be used to specify the changing behaviour of a model's state variables.

If implementors do decide to add scripting functionality to CellML documents, these scripts must be defined within elements placed in an application-specific extension namespace. Implementors are advised to follow the recommendations on the best practices for embedding and executing scripts described in Appendix B. The key recommendations are summarised in the following list:

- For interoperability, scripts should be defined using ECMAScript.
- The `<mathml:csymbol>` element should be used from within MathML markup to call scripts defined

using a non-MathML syntax. These elements must define a `definitionURL` that identifies the element containing the script, and an `encoding` attribute specifying the scripting language used.

- Function names (or the identifier used to reference a script) should be valid CellML identifiers, as defined in Section 2.2.1.
- The content of a `<mathml:csymbol>` element should be a human-readable identifier for the script, preferably the function name.
- Functions must be side-effect free. That is, a function must not assign values to variables that are not local to that function. In particular, functions must not alter the values of their arguments or global variables.

4.3 Examples

The CellML fragment in Figure 6 demonstrates how MathML can be employed within CellML to define mathematical expressions. This fragment is part of the definition of a component that represents the behaviour of the n gate from the potassium channel in the Hodgkin-Huxley squid axon model of 1952. The component contains two units definitions (with syntax defined in Section 5), two variable declarations (with syntax defined in Section 3), and a block of MathML that defines an expression calculating the α variable of the n gate as well as the rendering of this equation, which is given in Equation 2.

$$\alpha_n = 1.0 \frac{0.01(V + 10.0)}{\exp(0.1(V + 10.0)) - 1.0} \quad (2)$$

```
<component
  name="potassium_channel_n_gate"
  xmlns="http://www.cellml.org/cellml/1.0#"
  xmlns:cellml="http://www.cellml.org/cellml/1.0#"
  xmlns:cmeta="http://www.cellml.org/metadata/1.0#">

  <units name="per_millisecond">
    <unit prefix="milli" units="second" exponent="-1" />
  </units>
  <units name="millivolt">
    <unit prefix="milli" units="volt" />
  </units>
  <units name="per_millivolt">
    <unit prefix="milli" units="volt" exponent="-1" />
  </units>

  <variable name="alpha_n" units="per_millisecond" />
  <variable name="V" public_interface="in" units="millivolt" />
```

```

<math xmlns="http://www.w3.org/1998/Math/MathML">
  <semantics>
    <apply id="alpha_n_calculation"><eq />
      <ci> alpha_n </ci>
      <apply><times />
        <cn cellml:units="per_millisecond"> 1.0 </cn>
        <apply><divide />
          <apply><times />
            <cn cellml:units="per_millivolt"> 0.01 </cn>
            <apply><plus />
              <ci> V </ci>
              <cn cellml:units="millivolt"> 10.0 </cn>
            </apply>
          </apply>
          <apply><minus />
            <apply><exp />
              <apply><times />
                <cn cellml:units="per_millivolt"> 0.1 </cn>
                <apply><plus />
                  <ci> V </ci>
                  <cn cellml:units="millivolt"> 10.0 </cn>
                </apply>
              </apply>
            </apply>
          <cn cellml:units="dimensionless"> 1.0 </cn>
        </apply>
      </apply>
    </apply>
    <annotation-xml encoding="MathML-Presentation">
      <mrow>
        <mi> alpha_n </mi><mo> = </mo><mn> 1.0 </mn>
        <mfrac>
          <mrow>
            <mn> 0.01 </mn><mo> ( </mo>
              <mi> V </mi><mo> + </mo><mn> 10.0 </mn>
            <mo> ) </mo>
          </mrow>
          <mrow>
            <mo>exp</mo>

```

```

<mo> ( </mo><mn> 0.1 </mn><mo> ( </mo>
<mi> V </mi><mo> + </mo><mn> 10.0 </mn>
<mo> ) </mo><mo> ) </mo>
<mo> - </mo>
<mn> 1.0 </mn>
</mrow>
</mfrac>
</mrow>
</annotation-xml>
</semantics>
</math>
</component>

```

Figure 6 Part of the definition of a component that represents the behaviour of the n gate from the potassium channel in the Hodgkin-Huxley squid axon model of 1952. See text for more details.

Content that isn't defined using the MathML content markup elements can be associated with a MathML expression using the `<mathml:semantics>`, `<mathml:annotation>` and `<mathml:annotation-xml>` elements. The first child of a `<mathml:semantics>` element is the expression to be annotated, and the subsequent `<mathml:annotation>` and `<mathml:annotation-xml>` elements contain character data and XML annotations, respectively. In the CellML fragment in Figure 6, the expression of interest has been annotated with rendering information encoded using the MathML presentation markup elements. The MathML presentation elements are very flexible and it is possible to produce the same rendering of an equation in many ways - the choice of elements in Figure 6 is somewhat arbitrary.

The `<mathml:apply>` element at the top level of the expression defines an `id` attribute, which can be used to associate further metadata with the expression. The linking of metadata with elements in a CellML document is described in more detail in Section 8.2.

All of the `<mathml:cn>` elements in the equation define `cellml:units` attributes, which associate a units definition with the number delimited by the `<mathml:cn>` element. The inclusion of units in the equation allows CellML processing software to check that the dimensions of the terms in an equation are consistent, as discussed in Section 5. The presence of the unit scale factor on the right hand side of the equation is needed for the equation to have consistent dimensions.

4.4 Rules for CellML Documents

4.4.1 The `<mathml:math>` element

1. Allowed use of the `<mathml:math>` element

- The `<mathml:math>` element must only appear as a child of the following elements in the CellML

namespace: `<cellml:component>` and `<cellml:role>`. [In this and subsequent rules, the use of the `mathml` and `cellml` namespace prefixes indicates that elements and attributes are in the MathML and CellML namespaces, respectively. The `<mathml:math>` element may appear inside elements in the RDF and CellML Metadata namespaces if permitted by the relevant specifications, and may be used inside extension elements. When MathML elements occur within extension elements, CellML processing software is free to ignore them.]

- All elements in the MathML namespace that are within a `<mathml:math>` element, and not within a `<mathml:annotation>` or `<mathml:annotation-xml>` element, must be taken from the complete set of MathML content markup elements, as defined in Section 4.4 of the MathML 2.0 Recommendation, with the addition of the `<mathml:logbase>` element. [CellML only makes use of the content markup elements from MathML. However presentation markup elements may be used within the annotation elements to associate rendering information with expressions. The `<mathml:logbase>` element was erroneously omitted from the list of content markup elements in Section 4.4 of the MathML 2.0 Recommendation.]
- The content of a `<mathml:math>` element must conform to the MathML 2.0 Recommendation from the W3C.
- For interoperability, all elements in the MathML namespace that are within a `<mathml:math>` element, and not within a `<mathml:annotation>` or `<mathml:annotation-xml>` element should be taken from the CellML subset of MathML content markup elements defined in Figure 5. [The CellML subset is discussed further in Section 4.2.3. Note that this is an interoperability recommendation and not a firm rule.]

4.4.2 The `<mathml:ci>` element

1. Allowed use of the `<mathml:ci>` element

- After leading and trailing whitespace is removed, the content of a `<mathml:ci>` element must match the value of the name of a variable declared within the current component. [The `<mathml:ci>` element is used to reference variables from inside equations. Whitespace may be added before and/or after a variable's name to make the MathML more readable. The handling of whitespace in MathML is described in more detail in Section 2.4.6 of the MathML 2.0 Recommendation.]

4.4.3 The `<mathml:cn>` element

1. Allowed use of the `<mathml:cn>` element

- A `<mathml:cn>` element must define a `cellml:units` attribute. [All bare numbers in MathML content markup are enclosed in a `<mathml:cn>` element in the MathML namespace.]

2. Allowed values of the `cellml:units` attribute

- The value of the `cellml:units` attribute must be taken from the standard dictionary of units given in Section 5.2.1, or be the value of the `name` attribute on a `<cellml:units>` element defined in the current `<cellml:component>` or `<cellml:model>` element.

4.4.4 Modification of variables

- A mathematical expression defined using MathML must only modify the values of variables that belong to the current component. [Variables that belong to a component are those that are not declared with a `public_interface` or `private_interface` attribute value of "in".]

4.5 Rules for Processor Behaviour

4.5.1 Ordering of expressions

CellML processing software must not assume that the ordering of expressions within a CellML document has any significance.

4.5.2 Scope of expressions

CellML processing software must make no assumptions about the scope or domain of a mathematical expression defined within a model. Unless explicitly stated, all expressions hold for any and all combinations of independent variables.

4.5.3 Treatment of annotations

CellML processing software must assume that the content of the first child of a `<mathml:semantics>` element defines an expression describing the mathematical behaviour of the model. CellML processing software may ignore the content of `<mathml:annotation>` and `<mathml:annotation-xml>` elements.

5 Units

5.1 Introduction

One of the key features ensuring robustness and re-usability of CellML components and models is the requirement that units be associated with all variables and numbers in a CellML document. This allows components and models that declare variables with different units to be connected, as long as variables that are mapped to one another have the same dimensions. For instance, it is possible to map a variable declared with units of "pound/foot" to a variable declared with units of "kilogram/metre", but not to a variable declared with units of "mole/litre" or "kilogram-squared/metre". The explicit declaration of units also allows CellML processing software to check the consistency of each equation in a model.

5.2 Basic Structure

5.2.1 Dictionary of standard units

CellML provides a dictionary of standard units that may be used in variable declarations or attached to bare numbers in mathematics. References to these units should make use of the actual name of the units, rather than the standard abbreviation, thus avoiding confusion between units (e.g., metre) and prefixes (e.g., milli).

The full list of units that any CellML processing application is expected to recognise is given in Table 2. The keywords in the table comprise the SI base units, the SI derived units with special names and symbols, and some additional units commonly used in the types of biological models likely to be defined using CellML.

ampere	farad	katal	lux	pascal	tesla
becquerel	<i>gram</i>	kelvin	meter	radian	volt
candela	gray	kilogram	metre	second	watt
celsius	henry	<i>liter</i>	mole	siemens	weber
coulomb	hertz	<i>litre</i>	newton	sievert	
<i>dimensionless</i>	joule	lumen	ohm	steradian	

Table 2 The dictionary of units keywords that CellML processing applications are expected to recognise. Base SI units are printed in bold text, derived SI units are printed in plain text, and additions to the standard units defined purely for the convenience of model authors are italicised.

The SI base units are the foundation of the units system in CellML. The conversion of a variable's value between two sets of units involves the expansion of all units definitions to linear combinations of the SI base units and user-defined base units (described in Section 5.2.3). The list of SI base and derived units is taken from The International System of Units (SI), including the Year 2000 Supplement. The American spellings of **meter** and **liter** are taken from the NIST Guide for the Use of the International System of UNITS (SI). The SI standard defines the mathematical relationships between the SI derived units and the SI base units. These relationships are given in the right hand column of Table 3 in the Year 2000 Supplement, with the exception of **celsius**, which is related to **kelvin** as described in Section 2.1.1.5 of the SI standard.

The CellML units dictionary includes four non-SI units definitions for the convenience of modellers: *dimensionless*, *gram*, *liter* and *litre*. The only unfamiliar name on this list is *dimensionless*, which is used to indicate that a number or variable has no units associated with it. The mathematical relationships between *gram* and *litre* and the base SI units are given in Section 5.2.5.

5.2.2 User defined units

CellML also provides a facility whereby new units can be defined in terms of the units provided in the dictionary. This functionality allows the definition of units which are expressed as a scaled version of other units (as is the case for most imperial units), the definition of units which are made up of the product of other units, and even the creation of units that require an offset, such as degrees Fahrenheit. This allows model authors to work in whatever set of units they feel most comfortable, while still ensuring that their models can be integrated with those of other authors using other units.

New units are defined using the `<units>` element, which may be placed inside both `<model>` and `<component>` elements. When a `<units>` element is placed inside the `<model>` element, the units definition

may be referenced from within any component in the model. When a `<units>` element is placed inside a `<component>` element, the units definition may only be referenced from within that component.

Each `units` element must define a `name` attribute, which is used to reference the units definition elsewhere. The value of the `name` attribute must be unique across all `<units>` elements in the `<model>` or `<component>` element in which it is defined. If the value of the `name` attribute of a `<units>` element defined inside a `<component>` element matches the value of the `name` attribute on a `<units>` element defined inside the parent `<model>` element, then it will redefine the units, and all references to these units within the `<component>` element refer to the new definition. Model authors must not redefine any of the standard units. Therefore, the value of the `name` attribute must not equal one of the names from the standard units dictionary in Table 2.

A `<units>` element may also define a `base_units` attribute, the associated behaviour of which is discussed in Section 5.2.3. A `<units>` element may contain a set of `<unit>` elements that reference units from the dictionary or some previously defined units.

A `<unit>` element must not contain any elements in the CellML namespace, but may have up to five attributes. The `units` attribute is the only one that is required. It is used to set the base quantity for the current `<unit>` element, and its value must correspond to a keyword from the standard CellML units dictionary or to the value of the `name` attribute of a `<units>` element in the current component or model.

The definition of new units in terms of subunits may require the use of some combination of the optional `offset`, `prefix`, `exponent`, and `multiplier` attributes.

A `multiplier` attribute can be used to pre-multiply the quantity to be converted by any real scale factor. For instance, a multiplier of `0.45359237` is used to define a pound in terms of a kilogram. The `multiplier` attribute has a default value of `"1.0"`

The `offset` attribute is used to represent the addition of a constant in the transformation between the current units and the base units. This should only be necessary for the definition of temperature scales. For instance, an `offset` attribute value of `"32.0"` is needed to define Fahrenheit in terms of Celsius. The `offset` attribute has a default value of `"0.0"`.

The `prefix` attribute can be used to indicate a scale for the referenced units. It is included primarily for the convenience of modellers who want to define units that differ from another units definition only by an SI scale factor. Its value must be from the standard set of CellML prefix names given in Table 3 or be an integer, in which case the units are pre-multiplied by 10 to the power of this number. The default value of the `prefix` attribute is `"0"` (the referenced units are scaled by a factor of one).

namefactornamefactor

yotta	10 ²⁴	deci	10 ⁻¹
zetta	10 ²¹	centi	10 ⁻²

exa	10 ¹⁸	milli	10 ⁻³
peta	10 ¹⁵	micro	10 ⁻⁶
tera	10 ¹²	nano	10 ⁻⁹
giga	10 ⁹	pico	10 ⁻¹²
mega	10 ⁶	femto	10 ⁻¹⁵
kilo	10 ³	atto	10 ⁻¹⁸
hecto	10 ²	zepto	10 ⁻²¹
deka	10 ¹	yocto	10 ⁻²⁴

Table 3 The set of names that may be used in the `prefix` attribute on a `<unit>` element and the corresponding scale factors that will pre-multiply the unit.

The scale factor described by the `prefix` attribute and the units referenced by the `units` attribute are raised to a power equal to the value of the `exponent` attribute. The value of the `exponent` attribute must be a floating point number, and is typically an integer. The `exponent` attribute has a default value of "1.0". Note that an `exponent` attribute value of "0.0" has the effect of removing the parent `<unit>` element from the current units definition.

A *simple units* definition occurs when units are defined as a linear function of some previously defined simple units or base units. In a simple units definition, a `<units>` element contains only a single child `<unit>` element, that `<unit>` element has an `exponent` attribute value of "1.0", and the units definition referenced by the `units` attribute is one of the SI or user-defined base units or is itself a simple units definition. These are the only conditions under which a `<unit>` element may define an `offset` attribute with a value other than "0.0". The formula that expresses how the old units (referenced by the value of the `units` attribute on the `<unit>` element) are transformed into the new units (defined by the value of the `name` attribute on the parent `<units>` element) is given below.

$$x_{\text{new}}[\text{Units}] = (\text{multiplier prefix}) \left[\frac{\text{Units}}{\text{units}} \right] x_{\text{old}}[\text{units}] + \text{offset}[\text{Units}] \quad (3)$$

Terms in square brackets represent the units associated with values in the expression, which are italicised. x_{old} is the value to be transformed from the old units, and x_{new} is the resulting value in the new units. `Units` are the units being defined, and *multiplier*, *prefix*, `units` and *offset* correspond to the values of the appropriate attributes on the `<unit>` element.

Complex units are the product of multiple units. In a complex units definition, a `<units>` element contains more than one `<unit>` element or a `<unit>` element that defines an `exponent` attribute with a value other than "1.0". The conversion between the new units and the product of the constituent units is given by the formula below.

$$x_{\text{new}}[\text{Units}] = (m_1 \dots m_n p_1^{e_1} \dots p_n^{e_n}) \left[\frac{\text{Units}}{u_1^{e_1} \dots u_n^{e_n}} \right] x_{\text{old}}[u_1^{e_1} \dots u_n^{e_n}] \quad (4)$$

The m_i , p_i , u_i , and e_i terms refer to the values of the multiplier, prefix, units and exponent attributes on the i -th `<unit>` element respectively.

An `offset` attribute may not be defined on any `<unit>` elements that occur inside a complex units definition. When a complex units definition references a simple units definition, any offset associated with the simple units definition is removed. This means that conversions such as the one between degrees Fahrenheit per inch and degrees Celsius per centimetre involve only a scale factor.

5.2.3 New base units

A modeller might want to define and use units for which no simple conversion to SI units exist. A good example of this is pH, which is dimensionless, but uses a log scale. Ideally, pH should not simply be defined as dimensionless because software might then attempt to map variables defined with units of pH to any other dimensionless variables.

CellML addresses this by allowing the model author to indicate that a units definition is a new type of base unit, the definition of which cannot be resolved into simpler subunits. This is done by defining a `base_units` attribute value of "yes" on the `<units>` element. This element must then be left empty. The `base_units` attribute is optional and has a default value of "no". If the `base_units` attribute is omitted or assigned a value of "no", units are expected to be defined in terms of other units as described in Section 5.2.2.

The indiscriminate use of the `base_units` attribute is strongly discouraged, because it has a significant impact on the re-usability of models and components. In particular, the `base_units` attribute should not be used to restrict users to creating models with an application-specific dictionary of units, as this prevents the efficient exchange of CellML models with other applications.

Software that is checking the consistency of the units in an equation (described in more detail in Section 5.2.7) can stop the recursive resolution of units definitions when the only remaining units are base SI units and user-defined base units.

5.2.4 Expansion of units definitions

For interoperability, software that claims to perform units conversion when passing variables between components and/or claims to perform dimension consistency checking of equations should obtain results that are equivalent to those produced using the algorithms described in Appendix C.3.5 and Appendix C.3.6, respectively. Both of these algorithms make use of the algorithm defined in Appendix C.3.4 to fully expand units definitions into functions of the SI and user-defined base units.

For both simple and complex units definitions (as defined by Equation 3 and Equation 4, respectively), the algorithm recursively substitutes in equations expanding the unknown term x_{old} , stopping when the unknown term has only SI or user-defined base units.

Although this specification does not require software to implement this algorithm exactly, it is used extensively to demonstrate units conversion and dimension checking as described in Section 5.2.6 and Section 5.2.7,

respectively. Appendix C.4.2 provides examples of units definition expansion according to the algorithm described in Appendix C.3.4.

5.2.5 Expansion of the non-SI units definitions in the CellML dictionary

Having defined a mathematical notation in Section 5.2.2 and a technique for the expansion of units definitions, it is now possible to formally specify how the definitions of the non-SI units in Table 2 should be expanded. The CellML versions of these units definitions and the associated equations are given below. The definition of *liter* is identical to the definition of *litre*. As described in Section 5.2, *dimensionless* is not related to the SI units, and cannot be expanded.

```
<units name="gram">
  <unit multiplier="0.001" units="kilogram" />
</units>
```

$$x_{\text{new}}[\text{gram}] = 0.001 \left[\frac{\text{gram}}{\text{kilogram}} \right] x_{\text{old}}[\text{kilogram}] \quad (5)$$

```
<units name="litre">
  <unit multiplier="1000" prefix="centi" units="metre" exponent="3" />
</units>
```

$$\begin{aligned} x_{\text{new}}[\text{litre}] &= \left(1000 (10^{-2})^3 \right) \left[\frac{\text{litre}}{\text{metre}^3} \right] x_{\text{old}}[\text{metre}^3] \\ &= 0.001 \left[\frac{\text{litre}}{\text{metre}^3} \right] x_{\text{old}}[\text{metre}^3] \end{aligned} \quad (6)$$

5.2.6 Conversion between units definitions

Associating units definitions with every variable declaration in a component allows variables from components that make use of different sets of units to be mapped together, as long as the variables have the same dimensions. Appendix C.3.5 specifies a possible method for converting a numeric value from one set of units to another. CellML processing software is not required to be capable of converting between units definitions. However, for interoperability, software that does implement this functionality should achieve the same results as if this method were used, although the exact implementation may differ.

This implementation generates an expression that relates each units definition to SI and user-defined base units. This expression is obtained by recursively expanding each units definition as described in Appendix C.3.4, and then simplifying the result. The expression for the input units is then inverted to give an expression that relates the appropriate base units to the input units. This inverted expression is substituted into the expression for the target units, producing a single expression that relates the quantity to be converted from the input units to a corresponding quantity in the target units. The inversion and substitution process is demonstrated by example in Appendix C.4.3.

5.2.7 Equation dimension checking

The association of units with every variable and bare number that appears in an equation in a CellML document provides CellML processing software the opportunity to perform equation dimension checking. Verifying that equations have consistent dimensions can potentially catch many basic mathematical errors.

Appendix C.3.6 specifies a possible implementation of equation dimension checking. This implementation splits an equation into a tree of equation parts, in which each parent part is obtained by the application of a single operator to its children. The units definition on each leaf node (i.e., part without children) is expanded into base units, as described in Appendix C.3.4. The units definition for a node at a higher level of the tree is constructed by combining the units definition of its children. An equation has consistent dimensions if no errors are found while traversing the tree and if the fully expanded units definitions of the two nodes at the top level of the tree are equivalent, as defined in Appendix C.2.2.

CellML processing software is free to ignore units in mathematics and assume that equations are consistent. For interoperability, software that performs equation dimension checking should achieve the same results as if the implementation discussed in Appendix C.3.6 were used, although the exact implementation may differ.

This specification does not attempt to completely prevent model authors from creating invalid mathematics. Dimension consistency checking prevents modellers from adding variables with different dimensions but would not find errors in Equation 7 and Equation 8, which have different units but the same dimensions:

$$x[\text{volt}] = y[\text{volt}] + z[\text{millivolt}](7)$$

$$x[\text{inch}] = y[\text{metre}] + z[\text{nautical_mile}](8)$$

Although it would be technically possible (and useful) to find and correct such errors, CellML processing software is not required to be able to do so.

5.3 Examples

5.3.1 User-defined units and new base units

Figure 7 demonstrates how users can extend the set of units in the CellML dictionary by defining new sets of units.

```
<!-- User-defined Base Units -->
<units name="pH" base_units="yes" />

<!-- Simple Units Definitions -->
<units name="inch">
  <unit multiplier="2.54" prefix="centi" units="metre" />
</units>
```

```

<units name="fahrenheit">
<unit multiplier="1.8" units="celsius" offset="32.0" />
</units>

<!-- Complex Units Definitions -->
<units name="celsius_per_centimetre">
<unit units="celsius" />
<unit prefix="centi" units="metre" exponent="-1" />
</units>

<units name="fahrenheit_per_inch">
<unit units="fahrenheit" />
<unit units="inch" exponent="-1" />
</units>

<units name="pH_per_celsius">
<unit units="pH" />
<unit units="celsius" exponent="-1" />
</units>

```

Figure 7 Some examples of the use of the `<units>` element demonstrating the definition of simple and complex units.

5.3.2 Advanced examples

Examples of the expansions of units definitions, conversion between units definitions and equation dimension checking are given in Appendix C.4.

5.4 Rules for CellML Documents

Units are a fundamental part of a CellML model definition. In this section, formal rules are specified for the system of units definition introduced in Section 5.2.

5.4.1 The `<units>` element

1. Allowed use of the `<units>` element

- Both the `<model>` and `<component>` elements may contain any number of `<units>` elements.
- Each `<units>` element must define a `name` attribute, and may define a `base_units` attribute.
- If a `<units>` element defines a `base_units` attribute with a value of "yes", then that `<units>` element must contain only the following elements, which may appear in any order:
 - `<RDF>` elements in the RDF namespace.

- If a `<units>` element does not define a `base_units` attribute with a value of "yes", then that `<units>` element must contain only the following elements, which may appear in any order:
 - `<unit>` elements in the CellML namespace,
 - `<RDF>` elements in the RDF namespace.

2. Allowed values of the `name` attribute

- The value of the `name` attribute must be a valid CellML identifier as discussed in Section 2.2.1.
- The value of the `name` attribute must not equal one of the names defined in the standard dictionary of units in Table 2. [Model authors may not redefine the standard units.]
- The value of the `name` attribute must be unique across all `<units>` elements within a given `<model>` or `<component>` element. [Two `<units>` elements in the same `<model>` element may not have the same `name` attribute value, although a `<units>` element in a `<component>` element may share the same name as a `<units>` element in the parent `<model>` element. In this case, the units definition in the `<component>` element supercedes the model-wide definition when referenced inside that component.]

3. Allowed values of the `base_units` attribute

- If present, the value of the `base_units` attribute must be "yes" or "no".
- If not present, the value of the `base_units` attribute defaults to "no".

5.4.2 The `<unit>` element

1. Allowed use of the `<unit>` element

- A `<unit>` element must contain only the following elements:
 - `<RDF>` elements in the RDF namespace.
- Each `<unit>` element must define a `units` attribute. It may also define `prefix`, `exponent`, `multiplier`, and `offset` attributes.

2. Allowed values of the `units` attribute

- The value of the `units` attribute must be taken from the standard dictionary of units listed in Table 2 or be the value of the `name` attribute on a `<units>` element defined in the current `<component>` or `<model>` element.
- The value of the `units` attribute must not reference a units definition that contains `<unit>` elements that in turn directly or indirectly reference the current units definition. [This rule prevents circular units definitions. It must be possible to break down a complex units definition into SI and user-defined base units.]

3. Allowed values of the `prefix` attribute

- If present, the value of the `prefix` attribute must be an integer or a name taken from one of the name columns of Table 3. [The unit is scaled by 10 raised to the power of the specified integer or the factor corresponding to the specified name. Therefore, `prefix` attribute values of "centi" and "-2" are equivalent.]
- If not present, the value of the `prefix` attribute defaults to "0".

4. Allowed values of the `exponent` attribute

- If present, the value of the `exponent` attribute must be a real number.
- If not present, the value of the `exponent` attribute defaults to "1.0".

5. Allowed values of the `multiplier` attribute

- If present, the value of the `multiplier` attribute must be a real number.
- If not present, the value of the `multiplier` attribute defaults to "1.0".

6. Allowed values of the `offset` attribute

- If present, the value of the `offset` attribute must be a real number.
- If not present, the value of the `offset` attribute defaults to "0.0".

7. Proper use of the `offset` attribute

- A `<units>` element containing a `<unit>` element that defines an `offset` attribute with a value other than "0.0" must not contain other `<unit>` elements. [The `offset` attribute can only be used in a simple units definition, as defined in Section 5.2.2.]
- A `<unit>` element that defines an `offset` attribute with a value other than "0.0" must not define an `exponent` attribute with a value other than "1.0". [The `offset` attribute can only be used in a simple units definition, as defined in Section 5.2.2.]

5.5 Rules for Processor Behaviour

5.5.1 Resolving references to units definitions

The `<units>` element may be placed inside both `<model>` and `<component>` elements. When user-defined units are referenced by a variable or number declaration inside a component, the units definition is first looked for inside the current `<component>` element. If a matching units definition cannot be found, then the units definition is looked for in the parent `<model>` element.

5.5.2 Units associated with the MathML constants elements

This section defines the units associated with the MathML elements that appear in the **constants** subset of the CellML set defined in Section 4.2.3. These elements represent numerical values. Operators can be applied to combinations of these elements, variables and numbers in an equation. Units must be associated with these elements to allow for equation dimension checking.

The `<true>` and `<false>` elements have units of `cellml:boolean`, where `cellml:boolean` is a set of base units defined purely for use in this specification. (Note that users may not define their own `cellml:boolean` units, as this is not a valid CellML identifier.) `cellml:boolean` units are not associated with variables or numbers, but can be produced as the result of the application of relational or logical operators, as discussed in Appendix C.3.3.

The `<notanumber>`, `<pi>`, `<infinity>` and `<exponentiale>` elements all have units of dimensionless.

6 Grouping

6.1 Introduction

It is often useful to organise groups of components within a model into a hierarchical structure. This structure might reflect a logical organisation of components within the group or their physical configuration. CellML provides a single mechanism for the specification of both of these types of hierarchy that is based on a grouping scheme. This grouping scheme is general enough that it can be used within CellML documents to specify non-hierarchical grouping relationships between components.

In CellML, a *hierarchy* is a tree of components linked by parent-child relationships, where all of these relationships are of the same type. A hierarchy has a single root component and at least one child component. A model may contain numerous hierarchies of the same type. A component must only appear once within a set of hierarchies of a given type, but may appear in multiple hierarchies if each of these hierarchies is of a different type. CellML defines two types of relationship for use within the grouping scheme: *encapsulation* and *containment*.

The definition of a logical hierarchy of components in a network is known as "*encapsulation*". Encapsulation allows the modeller to hide part of a model by using a single component as an interface to a hidden submodel. The *parent* component hides the details of one or more *child* components from the rest of the model. Encapsulation provides a powerful mechanism for simplifying the structure of the model by preventing connections between specified sets of components. Components in the main model must not be connected to child components in the encapsulated submodel - all variables must be mapped through the encapsulating parent component. A component in the submodel must only be connected to its parent component, other components in the same submodel, and components that it encapsulates. A modeller wishing to re-use an encapsulated submodel may treat the submodel as a "black box", and deal exclusively with the interface presented by the encapsulating component.

The definition of physical hierarchies within a model is known as "*containment*". A model author can specify that one or more *child* components are physically inside of a *parent* component without describing the geometric aspects of the relationship in detail. This information would typically be used by CellML processing software to provide a physical representation of a model. A model may contain multiple types of containment hierarchy, which are differentiated based on names that the modeller assigns to these hierarchies.

Model authors are also free to extend the grouping scheme with user-defined types of relationships between components. These relationships need not be hierarchical in nature. However, CellML processing software is only required to recognise encapsulation and containment relationships.

Encapsulation and containment hierarchies do not add any mathematical information to the model. Model authors must not define their own grouping relationships that are intended to imply mathematical information.

Models may define multiple hierarchies of multiple types. CellML processing software is free to treat all hierarchies of the same type as separate hierarchies. Alternatively, it may combine all hierarchies of the same type into a single hierarchy by assuming that the root components of all explicitly defined hierarchies are children of a single *anonymous* component. This anonymous component is not explicitly defined within the

CellML document and has no properties.

6.2 Basic Structure

6.2.1 Definition of groups

Logical and physical hierarchies are both declared using the `<group>` element. This element must be a child of a `<model>` element. A `<group>` element can be used to define multiple hierarchies and associate multiple relationship types with each hierarchy. The definition of a hierarchy or set of hierarchies of the same type can be split up over multiple `<group>` elements, as long as all the children of a given parent component in a hierarchy appear in a single `<group>` element.

A `<group>` element must contain one or more `<relationship_ref>` elements, each of which must define a `relationship` attribute, the value of which references one type of relationship. CellML processing software is required to recognise two types of relationship: encapsulation and containment, which are indicated by `relationship` attribute values of "encapsulation" and "containment", respectively. Model authors can define new types of relationship by specifying a `relationship` attribute that is not in the CellML namespace on the `<relationship_ref>` element. All of the relationships referenced by the `<relationship_ref>` elements within a given `<group>` element are associated with all the parent-child pairs defined within that `<group>` element.

A `<group>` element must also contain one or more `<component_ref>` elements. Each `<component_ref>` element must define a `component` attribute, the value of which references a component within the current model. A parent-child link is created between components by nesting a `<component_ref>` element that references the child component inside a `<component_ref>` element that references the parent component. Multiple levels of nesting may be used within a single `<group>` element to define a hierarchy.

All `<component_ref>` elements defined immediately inside the `<group>` element must contain further `<component_ref>` elements when defining an encapsulation or containment hierarchy. This ensures that valid hierarchical structures are defined. Top-level `<component_ref>` elements need not contain further `<component_ref>` elements in `<group>` elements that reference only user-defined relationships. This allows the definition of non-hierarchical relationships.

A single hierarchy may be defined in multiple `<group>` elements. This occurs when a component is referenced in two groups that reference the same relationship type. However, all of the children of a given parent component must be defined within a single `<group>` element. Therefore, any given component can only be referenced once as a parent and once as a child for a given relationship type across the entire model. This simplifies the construction and validation of hierarchies.

A `<relationship_ref>` element may define a `name` attribute in addition to the required `relationship` attribute. The value of the `name` attribute on `<relationship_ref>` elements can be used to refine a given relationship type. This allows, for instance, the creation of several overlapping containment hierarchies within the same model, each with a different name. See Section 6.2.4 for more information on this.

Geometric containment relationship information is formally independent of logical encapsulation information, but CellML processing software is free to check for inconsistencies between the two relationships - it would generally be an error for an encapsulating component to be physically inside one of its encapsulated child components.

6.2.2 The *encapsulation* relationship type

Encapsulation allows the modeller to split a model into layers of complexity. A single component can be used to encapsulate a complex partial model, and thereby provide a unified interface for all information passing between that submodel and the rest of the model. This allows a modeller to refine the encapsulated submodel without having to make any changes to the rest of the model.

A model may contain any number of encapsulation hierarchies, as long as these do not overlap. If more than one hierarchy is explicitly defined, it may be useful to combine these into a single hierarchy by assigning all unencapsulated components an *anonymous* parent component. This anonymous component could make it easier to check that the hierarchies do not overlap and do not define any circular relationships between components.

The components in a model can be divided into four sets with respect to any given component (the *current* component). The set of all components immediately encapsulated by the current component is the *encapsulated set*. The *parent* component is the component that encapsulates the current component. Other components encapsulated by the same parent make up the *sibling set*. All other components, which are not available to make connections with the current component, make up the *hidden set*. If the current component is not encapsulated, then it has no parent and the sibling set consists of all other unencapsulated components in the model.

These sets are best demonstrated by example. Given the network shown in Figure 8, Table 4 lists the parent components and the components in the *encapsulated*, *sibling*, and *hidden* sets for a selected set of components picked as the *current* component.

network used to explain interface terminology

Figure 8 This simple model provides the basis for the demonstration of the concepts of *encapsulated sets*, *parents*, *sibling sets*, and *hidden sets*, as described in the text. The model consists of eight components each represented by a circle. The lines between the components represent connections, and a red arrowhead on one of these lines indicates that the component at the tail of the arrow is encapsulated by the component at the head of the arrow.

Current Component	Encapsulated Set	Parent	Sibling Set	Hidden Set
A	B, E	<i>anonymous</i>	G	C, D, F, H

B	C, D	A	E	F, G, H
C	<i>none</i>	B	D	A, E, F, G, H
E	F	A	B	C, D, G, H
G	H	<i>anonymous</i>	A	B, C, D, E, F

Table 4 This table lists the *parent* components, and the components in the *encapsulated*, *sibling*, and *hidden* sets for a selected few components from the example model in Figure 8. Components A and G are root components of separate hierarchies. It may be useful, however, to assign them an anonymous parent component that enables the formation of a single encapsulation hierarchy for the entire model.

Every variable must define its availability for use in other components. This is done with the `public_interface` and `private_interface` attributes on the `<variable>` element. The interface exposed to the parent component and components in the sibling set is defined by the `public_interface` attribute. The `private_interface` attribute defines the interface exposed to components in the encapsulated set. Each interface has three possible values: "in", "out", and "none", where "none" indicates the absence of an interface. The separation of interfaces allows the modeller to incrementally add complexity to an encapsulated submodel without changing the interface that the encapsulating component presents to the rest of the model.

The mappings allowed between variables declared in each component are defined by the public and private interfaces of each variable and the prohibition on connecting an encapsulated component to components other than its parent component, members of its sibling set, and any components it in turn encapsulates. Variables with a `public_interface` attribute value of "in" must be mapped to a single variable in a component in the sibling set with a `public_interface` attribute value of "out" or to a single variable in the parent of the current component with a `private_interface` attribute value of "out". Similarly, variables with a `public_interface` value of "out" may be mapped to variables in components in the sibling set with a `public_interface` attribute value of "in" or to variables in the parent component with a `private_interface` value of "in". Note that defining a `public_interface` attribute value of "out" on a variable makes it legal to map the variable to other variables, but does not require that such a mapping occur. If a variable has a `public_interface` attribute value of "none", it is neither input from or exposed to its parent or components in the sibling set.

Variables with a `private_interface` attribute value of "in" must be mapped to a single variable in a single component in the encapsulated set with a `public_interface` attribute value of "out". Variables with a `private_interface` attribute value of "out" may be mapped to variables in components in the encapsulated set with a `public_interface` attribute value of "in". If a variable has a `private_interface` attribute value of "none", it is neither input from or exposed to the components in the encapsulated set.

If either the `public_interface` attribute or the `private_interface` attribute of a variable have a value of "in", that variable is declared elsewhere and its value must not be mathematically modified in the current

component. Otherwise, the variable belongs to the current component. If both the `public_interface` and `private_interface` attributes of a variable have a value of "none", the variable can only be used in the current component and is invisible to all other components in the model.

The two interface attributes of a variable are completely independent with one exception: it is invalid for a variable to have both `public_interface` and `private_interface` attributes with a value of "in". An interface with a value of "in" reflects an unmet need in the current component that must be satisfied - this need can be met in either the public or private interface, but not both.

6.2.3 The *containment* relationship type

The containment relationship allows the modeller to specify that a particular component is physically inside another. This might be used by software to create a physical representation of the model. Containment relationships can be specified either in combination with or independent of encapsulation relationships. Containment relationships do not influence any aspect of model definition or behaviour.

6.2.4 Named containment relationship types

CellML allows the definition of multiple overlapping containment hierarchies in a given model. This functionality allows the modeller to define several different ways of organising a model, each of which might highlight a different aspect of the model's physical structure.

Multiple containment hierarchies are created by defining name attributes on the `<relationship_ref>` elements that have `relationship` attribute values of "containment". In effect, the introduction of a name attribute defines a new relationship type that has the same semantics as the unnamed containment relationship type. All containment hierarchies that share the same name are subject to the same rules that apply to any set of hierarchies that share the same relationship type. That is, each component must be referenced at most once as a parent or child for a given relationship type, and circular hierarchies must not be defined. Note that `<group>` elements that contain `<relationship_ref>` elements with a `relationship` attribute value of "containment" and that do not define a name attribute belong to a single relationship type that is separate from any named containment relationship types.

6.2.5 User-defined relationship types

Modellers are free to use the grouping syntax of CellML to organise model components in ways not defined by the CellML specification. To do this, the model author defines a new relationship type, the name of which is used as the value of the `relationship` attribute on the `<relationship_ref>` element. The `relationship` attribute must be placed in an extension namespace, because future versions of the CellML specification may define additional relationship types, the names of which could otherwise conflict with user-defined relationship types.

User-defined relationship types can be used in the definition of hierarchical relationships, and can also be used to define more generic *grouping* relationships. For example, a modeller may define a relationship type called `adjacency`, that indicates that any components referenced inside the group are physically adjacent to each

other.

Modellers are free to use the `name` attribute on the `<relationship_ref>` element to specify multiple hierarchies for user-defined relationship types, as is possible for the containment relationship type.

This specification does not provide a mechanism by which modellers may specify the meaning of a user-defined type of relationship. This definition must be provided by the processing software declaring the new relationship type.

6.3 Examples

Figure 9 demonstrates the use of the `<group>` element to define an encapsulation relationship. This example is taken from the two reaction pathway with encapsulation example from the examples section of the CellML website. It shows how a component representing an overall reaction (`total_reaction`) can encapsulate components representing intermediate reactions (`first_reaction` and `second_reaction`) and their by-products (C and D).

```
<group>
  <relationship_ref relationship="encapsulation" />
  <component_ref component="total_reaction">
    <component_ref component="first_reaction" />
    <component_ref component="second_reaction" />
    <component_ref component="C" />
    <component_ref component="D" />
  </component_ref>
</group>
```

Figure 9 Example demonstrating the use of the `<group>` element to define a logical encapsulation relationship. See text for more details.

Figure 10 demonstrates the use of the `<group>` element to define encapsulation and containment relationships, the construction of two named containment relationship types, and the specification of a custom relationship type (`adjacency`) in an extension namespace.

```
<group>
  <relationship_ref name="membrane" relationship="containment" />
  <component_ref component="cell">
    <component_ref component="cell_membrane" />
  </component_ref>
</group>
```

```

<group>
<relationship_ref relationship="encapsulation" />
<relationship_ref name="membrane" relationship="containment" />
<component_ref component="cell_membrane">
<component_ref component="sodium_channel" />
<component_ref component="calcium_channel" />
</component_ref>
</group>

<group>
<relationship_ref name="intracellular" relationship="containment" />
<component_ref component="cell">
<component_ref component="network_sarcoplasmic_reticulum" />
<component_ref component="junctional_sarcoplasmic_reticulum" />
</component_ref>
</group>

<group>
<relationship_ref
app:relationship="adjacency"
xmlns:app="http://www.software.com/cellml_processor" />
<component_ref component="network_sarcoplasmic_reticulum" />
<component_ref component="junctional_sarcoplasmic_reticulum" />
</group>

```

Figure 10 Examples demonstrating the use of the `<group>` element. See text for more details.

The first `<group>` element states that the `cell_membrane` component is physically inside the `cell` component, and that this particular containment relationship type is called `membrane`. The next `<group>` element states that the `sodium_channel` and `calcium_channel` components are both physically inside and logically encapsulated by the `cell_membrane` component. This containment relationship completes the `membrane` containment hierarchy. The encapsulation relationship prevents the sodium and calcium channel components from being connected to any components other than the `cell_membrane` component, each other, and any components they in turn encapsulate.

The third `<group>` element states that the two components representing parts of the sarcoplasmic reticulum are physically inside the cell, and that this relationship type is called `intracellular`. Finally, the fourth `<group>` element introduces the user-defined relationship `adjacency`, and states that the two sarcoplasmic reticulum components share this relationship. This relationship type is declared by putting the `relationship`

attribute in an extension namespace, and assigning it a value of "adjacency". Note that this relationship is not hierarchical in nature, and that CellML processing software is free to ignore the information provided by this group.

6.4 Rules for CellML Documents

6.4.1 The <group> element

1. Allowed use of the <group> element

- A <model> element may contain any number of <group> elements.
- A <group> element must contain only the following elements, which may appear in any order:
 - <relationship_ref> and <component_ref> elements in the CellML namespace,
 - <RDF> elements in the RDF namespace.

[Recommended practice is to define the CellML namespace child elements in a <group> element in the order stated above.]

- A <group> element must contain at least one <relationship_ref> element.
- A <group> element must contain at least one <component_ref> element.

6.4.2 The <relationship_ref> element

1. Allowed use of the <relationship_ref> element

- A <relationship_ref> element must contain only the following elements:
 - <RDF> elements in the RDF namespace.
- Each <relationship_ref> element must define a relationship attribute in either the CellML namespace or an extension namespace. It may also define a name attribute. [A relationship attribute declaring a user-defined relationship type is placed in an extension namespace. This restriction has been included to prevent conflicts with future versions of the CellML specification, which may define additional types of relationships in the CellML namespace.]

2. Allowed values of the relationship attribute

- The value of a relationship attribute in the CellML namespace must be "containment" or "encapsulation".

3. Allowed values of the name attribute

- The value of the name attribute must be a valid CellML identifier as discussed in Section 2.2.1. [Note that unlike most other name attributes, the value of the name attribute on a <relationship_ref> element is not expected to be unique across the current model. Instead, <group> elements that include <relationship_ref> elements that share the same name attribute value form the parts of a single hierarchy.]

4. Proper use of the name attribute

- A name attribute must not be defined on a <relationship_ref> element with a relationship attribute value of "encapsulation". [A model must define at most one

unnamed encapsulation hierarchy.]

5. Proper use of the **<relationship>** and **<name>** attributes

[The following rules together prevent the model author from referencing the same hierarchy more than once within a given **<group>** element.]

- A **<group>** element must not contain two or more **<relationship_ref>** elements that define a **relationship** attribute in a common namespace with the same value and that have the same **name** attribute value.
- A **<group>** element must not contain two or more **<relationship_ref>** elements that define a **relationship** attribute in a common namespace with the same value and do not define **name** attributes.

6.4.3 The **<component_ref>** element

1. Allowed use of the **<component_ref>** element

- A **<component_ref>** element must contain only the following elements, which may appear in any order:
 - **<component_ref>** elements in the CellML namespace,
 - **<RDF>** elements in the RDF namespace.
- A **<component_ref>** element must define a **component** attribute.

2. Proper use of the **<component_ref>** element

- A **<component_ref>** element that is defined immediately within a **<group>** element that contains a **<relationship_ref>** element with a **relationship** attribute value of "encapsulation" or "containment" must contain at least one child **<component_ref>** element. [Containment and encapsulation relationships must be hierarchical.]
- In a given hierarchy, only one of the **<component_ref>** elements that reference a given component may contain further **<component_ref>** elements. [This rule prevents a given component from being a parent more than once in a given hierarchy. A hierarchy is a set of components linked by a common type of parent-child relationship. The definition of a hierarchy may be split over multiple **<group>** elements, but the definition of a set of parent-child links must not be. It would be much more difficult to assemble a hierarchy from a CellML document if a set of parent-child links could be defined in multiple **<group>** elements.]
- In a given hierarchy, only one of the **<component_ref>** elements that reference a given component may be contained inside another **<component_ref>** element. [Complements the previous rule. This one prevents a given component from being a child more than once in a given hierarchy.]
- In a given hierarchy, a child component must not directly or indirectly contain its parent among its children. [A hierarchy must not be circular.]

3. Allowed values of the **component** attribute

- The value of the **component** attribute must equal the value of the **name** attribute of a **<component>** element contained within the current **<model>** element.

6.5 Rules for Processor Behaviour

6.5.1 Treatment of relationship types in a single model

A given relationship type must have the same semantics across a model and at all levels in every hierarchy associated with that relationship type. The semantics of the encapsulation and containment relationship types are defined in Section 6.2.2 and Section 6.2.3, respectively.

Within a given `<model>` element, any hierarchies defined in `<group>` elements that contain `<relationship_ref>` elements with identical `relationship` and `name` attribute values belong to the same relationship type, and must be treated as such. Any hierarchies defined in `<group>` elements that contain `<relationship_ref>` elements with identical `relationship` attribute values and undefined `name` attributes belong to the same relationship type, and must be treated as such.

6.5.2 Groups must not imply mathematical information

Modellers must not use CellML groups to add mathematical information to the model. Modellers must not define their own types of relationships that imply mathematical behaviour. This ensures that the mathematical behaviour of a model can be properly reproduced by all CellML processing software.

6.5.3 Groups must not imply metadata information

Modellers must not use CellML groups to associate properties or classification information with sets of components. The metadata functionality is the proper method for making such associations. This increases the chance of that information being used by a range of CellML processing software.

7 Reactions

7.1 Introduction

CellML is intended to be used to represent many different types of models. Therefore, its basic structure is general. Models are primarily specified by explicitly defining mathematics using MathML. It is possible to specify a model purely in terms of mathematics, without using any of the elements defined in this section of the specification. However, in some types of models, information is lost when reducing the model to pure mathematics. For instance, in biochemical pathway models it will not always be straightforward, or even possible, to unambiguously determine from the mathematical rate laws which variables represent inhibitors or activators in the reactions. Therefore, some additional elements are needed in CellML to fully capture the information in biochemical pathway models.

7.1.1 Pathway model representations supported by CellML

Three fundamental representations of reaction/pathway models must be supported by CellML:

- **Mathematical Equations:** These are any valid mathematical equations that describe the model. For

example, they may be ordinary differential equations that define kinetic reaction rate laws and the rate of change of the concentration of species participating in the modelled reactions.

- **Chemical Expressions:** These are the stoichiometric expressions (such as $A + B \leftrightarrow 2C + D$) used by chemists to represent reactions.
- **Pathway Diagrams:** These are the stylised drawings commonly used by biochemists and cell biologists to represent interactions among participants in reactions. Some examples of pathway diagrams are shown in Section 7.3.

It is important that CellML be able to store the information needed to reproduce unambiguously any of these representations of a model. It is also important to minimise duplication of information within the model definition, because duplication can lead to inconsistencies. Therefore, CellML integrates the information needed to support the three types of model representation.

The integration process has resulted in the introduction of a CellML syntax that implies a mathematical relationship between variables in the current component. In this section of the specification, *explicit* mathematics refers to equations defined using MathML, and *implicit* mathematics refers to equations implied from the CellML syntax.

7.1.2 Qualitative and quantitative pathway models

CellML supports both quantitative and qualitative pathway models. Many types of models are commonly referred to as "qualitative". Some of these are mathematically specified, while others are not. For the purposes of this specification, qualitative pathway models consist solely of information about how the different chemical species in the pathway relate, and contain no mathematics. However, the stoichiometry of the reactions may be known. In other words, there is no mathematical representation of the model, but there may still be a pathway diagram and chemical expressions that represent the model. Because a qualitative model has no mathematics, CellML processing software is not required to be able to run a simulation using a qualitative model. However, some software may support simple simulations using such models.

Any model in which the change of concentration of a chemical species participating in a reaction is implicitly or explicitly defined is a *quantitative pathway model*. All others are *qualitative pathway models*.

7.2 Basic Structure

The `<reaction>` element is used to store information associated with a single reaction. It may only appear inside a `<component>` element. Examples demonstrating the use of the `<reaction>` element are presented in Section 7.3. It is possible for a single `<component>` element to contain more than one `<reaction>` element. However, this makes it difficult to re-use the individual reactions, and is therefore not the recommended best practice. The `<reaction>` element may define a `reversible` attribute, the value of which indicates whether or not the reaction is reversible. The default value of the `reversible` attribute is "yes".

The reaction element contains multiple `<variable_ref>` elements, each of which references a variable that

participates in the reaction. The recommended best practice is to create a `<variable_ref>` element for each variable representing the concentration of a chemical species that participates in a reaction, as well as one for the variable representing the rate of the reaction. The required `variable` attribute is the only attribute on the `<variable_ref>` element. Its value is the name of the referenced variable. This variable must be declared in the current `<component>` element.

Each `<variable_ref>` element contains one or more `<role>` elements. A `<role>` element must not contain any elements in the CellML namespace, but may have up to four attributes. The required `role` attribute specifies the way in which the variable participates in the reaction. The `role` attribute must have a value of "reactant", "product", "catalyst", "activator", "inhibitor", "modifier", or "rate". The meaning associated with each value is defined in Section 7.4. The optional `direction` attribute may be used on `<role>` elements in reversible reactions. If defined, it must have a value of "forward", "reverse", or "both". Its value indicates the direction of the reaction for which the role is relevant. It has a default value of "forward". The optional `delta_variable` attribute indicates which variable is used to store the change in concentration of the species represented by the variable referenced by the current `<variable_ref>` element. The optional `stoichiometry` attribute stores the stoichiometry of the current variable relative to the other reaction participants. Section 7.4 contains detailed rules for the use of these attributes.

The `<role>` elements may also contain `<math>` elements in the MathML namespace, which define equations using MathML. Although it is not required, it is recommended best practice to store all of the equations that relate to a reaction inside the appropriate `<role>` elements in the `<reaction>` element. This makes the `<reaction>` element more re-usable. In addition, defining mathematics inside a `<role>` element has the effect of associating the equations with the variable referenced by the containing `<variable_ref>` element, in the role defined by the `<role>` element. This enables CellML processing software to present the equations in a more meaningful context. For instance, it may group all of the relationships between the rate variable and the delta variables for all of the reactants and products, or it may display these equations in a different color.

There are three uses for equations inside `<role>` elements:

- If the `role` attribute value is "rate", any enclosed equations calculate the kinetic rate law (i.e., calculate the value of the referenced variable) and the value of intermediate variables used in the rate law equation.
- If the `role` attribute value is "reactant" or "product", the equations calculate the relationship between the general reaction rate and the rate of change of the species represented by the referenced variable (i.e., calculate the value of the variable named in the `delta_variable` attribute), and calculate any intermediate variables used in this relationship.
- In all other cases, the equations relate an intermediate variable used in the rate calculation to the variable referenced by the containing `<variable_ref>` element. For instance, it would be appropriate to calculate an effective concentration of a catalyst inside the `<role>` element contained by the `<variable_ref>` element that references the variable representing the actual concentration of the catalyst.

CellML processing software is not required to be able to deduce the stoichiometry of a reaction from explicit mathematics. Therefore, it is strongly recommended that the `stoichiometry` and `delta_variable` attributes be used instead of explicit mathematics if the concentration change is simply the reaction rate multiplied by the stoichiometry. (The rules for deriving this mathematical relationship from the `stoichiometry` attribute are defined in Section 7.5.5.)

7.3 Examples

This section contains two examples demonstrating the recommended use of the `<reaction>`, `<variable_ref>` and `<role>` elements to define two basic reactions. The mathematics defining the reaction rate have been omitted in these examples. See the signal transduction model examples section of the CellML website for further examples.

Figure 11 shows a pathway diagram representation of the following reversible reaction:




Figure 12 demonstrates the use of CellML to define this reaction. There are five `<variable_ref>` elements in the `<reaction>` element: one for each variable representing the concentration of a chemical species participating in the reaction, and one for the variable representing the general reaction rate. Note that the `stoichiometry` attribute has a value of "2" for the variable representing the chemical species C, since this species appears with a stoichiometry of 2 in the chemical expression. The `reversible` attribute on the `<reaction>` element and the `direction` attributes on the `<variable_ref>` elements have their default values ("yes" and "forward", respectively) and therefore could have been omitted, but they are included for clarity.

reaction picture

Figure 11 A typical pathway diagram representation of the simple reversible reaction $A + B \rightleftharpoons 2C + D$.

```
<reaction reversible="yes">
  <variable_ref variable="A">
    <role
      role="reactant" direction="forward"
      delta_variable="delta_A" stoichiometry="1" />
    </variable_ref>

  <variable_ref variable="B">
    <role
      role="reactant" direction="forward"
```

```

delta_variable="delta_B" stoichiometry="1" />
</variable_ref>

<variable_ref variable="C">
<role
role="product" direction="forward"
delta_variable="delta_C" stoichiometry="2" />
</variable_ref>

<variable_ref variable="D">
<role
role="product" direction="forward"
delta_variable="delta_D" stoichiometry="1" />
</variable_ref>

<variable_ref variable="r">
<role role="rate">
<math xmlns="http://www.w3.org/1998/Math/MathML">
... <!-- reaction rate math -->
</math>
</role>
</variable_ref>
</reaction>

```

Figure 12 The CellML definition of the simple reversible reaction $A + B \rightleftharpoons 2C + D$. See text for more details.

Figure 13 shows the pathway diagram for the following irreversible, catalyzed reaction, which exhibits product-inhibition:



The CellML definition of this reaction is shown in Figure 14.

reaction picture

Figure 13 A typical pathway diagram representation of the irreversible reaction $A + B \rightarrow D$ (catalyzed by C, inhibited by D).

```

<reaction reversible="no">
  <variable_ref variable="A">
    <role role="reactant" delta_variable="delta_A" stoichiometry="1" />
  </variable_ref>

  <variable_ref variable="B">
    <role role="reactant" delta_variable="delta_B" stoichiometry="1" />
  </variable_ref>

  <variable_ref variable="C">
    <role role="catalyst" />
  </variable_ref>

  <variable_ref variable="D">
    <role role="product" delta_variable="delta_D" stoichiometry="1" />
    <role role="inhibitor" stoichiometry="1" />
  </variable_ref>

  <variable_ref variable="r">
    <role role="rate">
      <math xmlns="http://www.w3.org/1998/Math/MathML">
        ... <!-- reaction rate math -->
      </math>
    </role>
  </variable_ref>

</reaction>

```

Figure 14 The CellML definition of the irreversible reaction $A + B \rightarrow D$ (catalyzed by C, inhibited by D). See text for more details.

The `<variable_ref>` element that references the variable representing the concentration of species D now contains two `<role>` elements, one with information about D as a product and the other with information about D as an inhibitor. In this example, D has the same stoichiometry in both roles, but this would not necessarily need to be the case.

7.4 Rules for CellML Documents

7.4.1 The `<reaction>` element

1. Allowed use of the <reaction> element

- A <component> element may contain any number of <reaction> elements. [The use of multiple <reaction> elements within a single <component> element is discouraged.]
- A <reaction> element must contain only the following elements, which may appear in any order:
 - <variable_ref> elements in the CellML namespace,
 - <RDF> elements in the RDF namespace.

[The recommended best practice is to define one <variable_ref> element for each variable representing a chemical species that participates in the reaction, and one <variable_ref> element for the variable representing the rate of the reaction.]

- Each <reaction> element must contain at least one <variable_ref> element.
- The <reaction> element may define a reversible attribute.

2. Allowed values of the reversible attribute

- If present, the reversible attribute must have a value of "yes" or "no".
- If not present, the value of the reversible attribute defaults to "yes". [It is recommended to always explicitly define the value of this attribute.]

3. Proper use of the <reaction> element in encapsulating components

[It is often convenient to include a <reaction> element in a component that is encapsulating several intermediate reactions (see Section 6 for more information about encapsulation). The encapsulating component represents an overall, or total, reaction, which can be represented by a <reaction> element. This total reaction is effectively qualitative, because the mathematics representing the progression of the total reaction are defined in the components representing the intermediate reactions.]

- A <reaction> element in an encapsulating component must not contain delta_variable attributes on the <role> elements or explicit mathematics defining the overall reaction rate or the changes in concentration of the species that participate in the total reaction. [A CellML document should not define an inconsistent set of equations. This rule prevents authors inadvertently introducing explicit or implicit mathematics in an encapsulating component that duplicates or contradicts mathematics (either explicit or implicit) defined in the encapsulated components.]

7.4.2 The <variable_ref> element

1. Allowed use of the <variable_ref> element

- A <variable_ref> element must contain only the following elements, which may appear in any order:
 - <role> elements in the CellML namespace,
 - <RDF> elements in the RDF namespace.
- Each <variable_ref> element must contain at least one <role> element. [The recommended best practice is to define one <role> element for each role assumed by the chemical species represented by the referenced variable.]
- Each <variable_ref> element must define a variable attribute.

2. Allowed values of the variable attribute

- The value of the `variable` attribute on a `<variable_ref>` element within a `<reaction>` element must equal the value of the `name` attribute on a `<variable>` element defined inside the current `<component>` element.
- The value of the `variable` attribute must be unique across all `<variable_ref>` elements contained within the parent `<reaction>` element. [A variable must only be referenced once in a single reaction.]

7.4.3 The `<role>` element

1. Allowed use of the `<role>` element

- A `<role>` element must contain only the following elements, which may appear in any order:
 - `<math>` elements in the MathML namespace,
 - `<RDF>` elements in the RDF namespace.

[Some rules for the use of mathematics in `<role>` elements are provided below, and rules for the `<math>` element and its children are given in Section 4.]

- Each `<role>` element must define a `role` attribute. It may also define `delta_variable`, `direction`, and `stoichiometry` attributes.

2. Allowed values of the `role` attribute

- The `role` attribute must take one of the following seven values:
 - `"reactant"`: the species represented by the referenced variable is one of the species consumed or transformed by the reaction (in the forward direction). Reactants are also often called substrates.
 - `"product"`: the species represented by the referenced variable is one of the species produced by the reaction (in the forward direction).
 - `"catalyst"`: the species represented by the referenced variable catalyzes the reaction. In biochemical pathways such a species will almost always be an enzyme and will almost always occur with a `stoichiometry` attribute value of `"1"`.
 - `"activator"`: the species represented by the referenced variable enhances the reaction. Activators can occur with any stoichiometry. An activator will usually be a small molecule that increases the activity of an enzyme catalyzing the reaction. However, the detailed reaction representing this activation of the enzyme may not be included in the model. Instead, the activator may be represented as directly affecting the kinetics of the catalyzed reaction.
 - `"inhibitor"`: the species represented by the referenced variable inhibits the reaction. Inhibitors can occur with any stoichiometry. An inhibitor will usually be a species that inhibits the activity of an enzyme catalyzing the reaction. However, the detailed reaction representing this inhibition of the enzyme may not be included in the model. Instead, the inhibitor may be represented as directly affecting the kinetics of the catalyzed reaction.
 - `"modifier"`: the species represented by the referenced variable modifies the reaction in some unspecified way.
 - `"rate"`: the referenced variable represents the rate of the reaction.

3. Proper use of the role attribute

- A `<reaction>` element must contain no more than one `<variable_ref>` element with a `<role>` element with a `role` attribute with a value of "rate". [There may only be one rate variable per reaction.]
- A `<variable_ref>` element that contains a `<role>` element with a `role` attribute value of "rate" must not contain other `<role>` elements. [The variable assigned the "rate" role may not be assigned any other roles.]
- A `<role>` element with a `role` attribute value of "rate" must not define `delta_variable` or `stoichiometry` attributes. [The `delta_variable` and `stoichiometry` attributes have no meaning for a rate variable.]

4. Allowed values of the direction attribute

- If present, the `direction` attribute must take one of the following three values:
 - "forward": the value of the `role` attribute is the role of the referenced variable in the reaction when running in the "favoured" direction. The favoured direction is the one in which the the reactants are being consumed (i.e., the time-derivatives of their concentrations are negative), as defined by the kinetic rate law.
 - "reverse": the value of the `role` attribute is the role of the referenced variable in the reaction when running opposite to the "favoured" direction. In this direction, the reactants (as defined by the kinetic rate law) are being produced.
 - "both": the value of the `role` attribute is the role of the referenced variable in both directions of the reaction.
- If not present, the value of the `direction` attribute defaults to "forward".

5. Proper use of the direction attribute

- A `direction` attribute on a `<role>` element that is inside a `<reaction>` element with a `reversible` attribute value of "no" must have a value of "forward". [Only reversible reactions may occur in two directions.]
- A `direction` attribute on a `<role>` element for which the `role` attribute has a value of "rate", "reactant" or "product" must have a value of "forward". [Variables representing the reaction rate, a reactant or a product should always be labelled as such in the forward direction. To do otherwise would cause the implicit mathematics defined by the `delta_variable` and `stoichiometry` attributes on the reactant and product variables to be erroneous.]
- Each `<role>` element contained in a given `<variable_ref>` element must have a unique combination of values for the `role` and `direction` attributes. [Defining two `<role>` elements with the same `role` and `direction` attribute values would allow the definition of inconsistent stoichiometries or multiple delta variables for a single variable. Both of these situations would imply inconsistent mathematics.]

6. Allowed values of the stoichiometry attribute

- If present, the value of the `stoichiometry` attribute must be a real number. [In most cases, the value will be an integer. However, a valid CellML model may use fractional stoichiometries.]

7. Allowed values of the delta_variable attribute

- If present, the value of the `delta_variable` attribute must equal the name attribute on a `<variable>` element defined inside the current `<component>` element.
- If present, the value of the `delta_variable` attribute must be unique across all `<role>` elements contained within the parent `<component>` element. [One variable cannot represent the rate of change in concentration of more than one species. The value of the `delta_variable` attribute must be unique across the entire `<component>` element because it is legal (but not recommended) to include more than one `<reaction>` element in a single component.]

8. Proper use of the `delta_variable` attribute

- A `delta_variable` attribute must only appear on `<role>` elements in which the `role` attribute has a value of "reactant" or "product". [It is only in these roles that a chemical species may undergo a change in concentration.]
- A `<role>` element on which a `delta_variable` attribute is declared must also either declare a `stoichiometry` attribute or contain at least one `<math>` element in the MathML namespace. [The combination of the `delta_variable` attribute and the `stoichiometry` attribute implies a mathematical relationship between the variable referenced in the `delta_variable` attribute and the variable assigned the role of "rate", as defined in Section 7.5.5. If the `stoichiometry` attribute is absent, the relationship between the variable assigned the role of "rate" and the variable named in the `delta_variable` attribute must be defined using MathML.]
- A `<role>` element on which the `stoichiometry` and `delta_variable` attributes are both defined must not contain `<math>` elements in the MathML namespace. [The equations in a `<math>` element inside a `<role>` element for which the `role` attribute is "reactant" or "product" must relate the variable named in the `delta_variable` attribute to the variable assigned the role of "rate". Such equations would contradict the relationship implied by the `delta_variable` and `stoichiometry` attributes, as defined in Section 7.5.5.]
- If the `delta_variable` and `stoichiometry` attributes are both declared on any single reaction participant, a `<variable_ref>` element must be provided for the variable that represents the reaction rate. This `<variable_ref>` must contain exactly one `<role>` element, with a `role` attribute equal to "rate". [The reverse is not true: a variable may be assigned a role of "rate" even if the "reactant" and "product" variables do not define `delta_variable` attributes. In this case, the modeller may choose to provide explicit mathematics relating the "rate" variable to the change in concentration of the various chemical species.]

9. Proper use of a `<math>` element inside a `<role>` element

- A `<math>` element in the MathML namespace inside a `<role>` element must define equations that are relevant to the variable referenced by the containing `<variable_ref>` element, acting in the role defined by the `role` attribute on the `<role>` element. [The meaning of "relevant" in this context is discussed in Section 7.5.6.]

7.5 Rules for Processor Behaviour

7.5.1 Implications of the `reversible` attribute

If the `reversible` attribute on a `<reaction>` element has a value of "yes", it is assumed that all reactants in the forward direction are products in the reverse direction and vice versa. Similarly, all products in the forward direction are assumed to be reactants in the reverse direction and vice versa.

7.5.2 The absence of a stoichiometry attribute

CellML processing software must not make any assumptions if a `stoichiometry` attribute is not defined on a `<role>` element. The absence of a stoichiometry value specifically does **not** imply a stoichiometry of "1". Instead, it would usually mean that the stoichiometry of the reaction is unknown.

7.5.3 Chemical information implied by the stoichiometry attribute

The value of the `stoichiometry` attribute on a `<role>` element is defined to be the stoichiometry of the chemical species whose concentration is represented by the variable referenced by the containing `<variable_ref>` element. This stoichiometry can be used to produce the chemical expression representation of the model.

7.5.4 The absence of a delta_variable attribute

CellML processing software must not make any assumptions if a `delta_variable` attribute is not defined on a `<role>` element.

7.5.5 Math implied by the delta_variable and stoichiometry attributes

The use of the `delta_variable` and `stoichiometry` attributes on a `<role>` element implies the following mathematical relationship between the declared delta variable and the rate variable (where the variable representing the reaction rate will have a negative value when the reaction is proceeding in the forward direction):

- For reactants: $\text{delta_variable} = (\text{stoichiometry})(\text{rate})$
- For products: $\text{delta_variable} = -(\text{stoichiometry})(\text{rate})$

The two reactions shown in Figure 15 are mathematically equivalent. The representation in the first reaction in Figure 15 is the recommended best practice, because processing applications are not required to be able to extract the stoichiometry from an explicit MathML definition such as the one shown in the second reaction.

Explicit mathematics should only be used in cases where the implicit formulation would be inappropriate. Some examples of such cases are:

- If the stoichiometry of a reaction is unknown, but the modeller still wishes to relate the rate of change of a particular chemical species to the general reaction rate. Defining the `stoichiometry` attribute implies that the stoichiometry is known to equal the value of that attribute.
- If the modeller wishes to experiment with the stoichiometry of a species in different simulations using

the model. (In this case, it might be easier if the stoichiometry is defined as a variable.)

- If the math implied from the recommended formulation would be incorrect, i.e., in the rare cases when a more complex function is needed to relate the change in concentration of a species to the reaction rate.

In all of these cases, it is recommended best practice to put the mathematical expression used to define the change in concentration of a species inside the `<role>` element contained in the `<variable_ref>` element referring to the variable representing the concentration of that species.

```
<!--
```

```
The recommended best practice for calculating the value of delta_A, which
is an implied function of stoichiometry and reaction rate.
```

```
-->
```

```
<reaction reversible="yes">
<variable_ref variable="A">
<role
role="reactant" direction="forward"
delta_variable="delta_A" stoichiometry="2" />
</variable_ref>
```

```
<variable_ref variable="r">
<role role="rate">
<math xmlns="http://www.w3.org/1998/Math/MathML">
... <!-- reaction rate math -->
</math>
</role>
</variable_ref>
</reaction>
```

```
<!--
```

```
In this reaction, the value of delta_A is calculated explicitly using MathML.
This is not the recommended best practice.
```

```
-->
```

```
<reaction reversible="yes">
<variable_ref variable="A">
<role
role="reactant" direction="forward"
delta_variable="delta_A">
<math xmlns="http://www.w3.org/1998/Math/MathML">
<apply><eq />
```

```

<ci> delta_A </ci>
<apply><times />
<cn cellml:units="dimensionless"> 2.0 </cn>
<ci> r </ci>
</apply>
</math>
</role>
</variable_ref>

<variable_ref variable="r">
<role role="rate">
<math xmlns="http://www.w3.org/1998/Math/MathML">
... <!-- reaction rate math -->
</math>
</role>
</variable_ref>
</reaction>

```

Figure 15 The top `<reaction>` element shows the recommended best practice for defining the change in concentration `delta_A` of a chemical species A with respect to the reaction rate `r`. The second `<reaction>` element shows an equivalent representation using an explicit MathML definition. Use of this formulation is not recommended. The MathML blocks defining the rate laws are omitted.

It is an error to explicitly declare mathematics that conflicts with or duplicates implied mathematics. Therefore, a modeller must not declare a `stoichiometry` attribute and `delta_variable` attribute in addition to explicit math relating the change in concentration of the referenced species to the reaction rate.

7.5.6 Meaning of mathematics in reactions

Equations defined in `<math>` elements in the MathML namespace inside a `<role>` element must be relevant to the the variable referenced by the parent `<variable_ref>` element, acting in the role defined by the value of the `role` attribute. This means that:

- If the `role` attribute value is "rate", the equations must calculate the kinetic rate law (i.e., calculate the value of the referenced variable). Intermediate calculations related to the calculation of the rate are also allowed. Conventionally, the variable representing the reaction rate will have a negative value when the reaction is proceeding in the forward direction.
- If the `role` attribute value is "reactant" or "product", the equations must calculate the relationship between the general reaction rate and the rate of change of the species represented by the referenced

variable (i.e., calculate the value of variable named in the `delta_variable` attribute). Intermediate calculations related to the calculation of the delta variable are also allowed.

- In all other cases, the equations must relate an intermediate variable used in the rate calculation to the variable referenced by the containing `<variable_ref>` element. For example, it would be appropriate to calculate an effective concentration of an inhibitor or catalyst in the `<role>` element contained in the `<variable_ref>` element that references the variable representing the actual concentration of that species.

7.5.7 Resolution of inconsistencies

Duplication of information is avoided as much as possible. However, because modellers must be free to define arbitrary rate laws, it was not possible to eliminate all information duplication. For instance, CellML processing software is not expected to be able to deduce all information about a reaction from kinetic laws of arbitrary form, even though most information is in fact represented in these laws. Therefore, there is a possibility that the information in the mathematics and the information in the `<reaction>` element may be inconsistent.

It is anticipated that most modellers will define CellML models using some sort of processing software, which can reasonably be expected to write consistent mathematics. However, since CellML is a text-based format, modellers may also create or edit models by hand, and in doing so risk creating inconsistent models.

The following rules govern the required behaviour of CellML conformant processing software in the event that information in the explicitly and implicitly defined mathematics do not agree:

- CellML processing software may check for inconsistencies between mathematics explicitly defined using MathML and mathematics implicitly defined using the attributes on the `<role>` element. If software does check, it is recommended that it notify the user if inconsistencies are detected.
- If inconsistencies are found, CellML processing software must use mathematics explicitly defined using MathML when **running** a simulation with the model.
- CellML processing software may treat inconsistencies as it chooses when **representing** the model. For instance, software may ignore inconsistencies when rendering pathway diagrams or chemical expressions.

8 Metadata Framework

8.1 Introduction

Metadata is "*data about data*". In a CellML document, the principal data is the structure and mathematics of a biological model. Information that provides context for this data is metadata. Metadata can be included in a CellML document to facilitate searches of collections of models and model components. It provides a means for a modeller to include structured descriptive information about the model, which can help other modellers determine whether they can incorporate the model into their own work.

This section of the CellML specification presents a framework for the use of metadata in a CellML document.

Methods for identifying types of metadata within that framework are recommended in the CellML Metadata Specification. The use of these methods ensures reliable extraction of metadata from CellML documents across all processors aware of CellML Metadata. The CellML Metadata specification is being developed independently of the CellML specification.

All metadata is optional. A model without any metadata is a valid CellML model. However, it is recommended that a CellML document author provide as much metadata as possible, particularly his/her name and contact information and a reference for a paper that describes the development of the model.

8.2 Basic Structure

Metadata should be embedded in a CellML document using the Resource Description Framework (RDF), the syntax of which is defined in the RDF Model and Syntax Recommendation. For interoperability, CellML processing software should make use of the methods for identifying types of metadata outlined in the CellML Metadata Specification.

Section 2.2.2 defines two metadata namespaces that CellML processing software is expected to recognise, and recommended prefixes to which these namespaces should be mapped. RDF elements are placed in the RDF namespace, which should be mapped to the prefix `rdf`. CellML Metadata elements and attributes have their own namespace which should be mapped to the prefix `cmeta`.

CellML processing software is free to ignore any and all metadata. However, it is recommended that software at least display metadata. Model authors are free to develop their own RDF schema for metadata, or to store metadata in another format by using the CellML extension mechanism described in Section 2.2.3. However, doing so decreases the likelihood that CellML processing software will be able to do anything useful with the metadata in a CellML document.

Metadata is defined within an `<rdf:RDF>` element as shown in Figure 16. The `rdf`, `cellml`, and `cmeta` prefixes are used throughout this section to indicate that elements and attributes are in the RDF, CellML and CellML Metadata namespaces, respectively. The recommended best practice is to define the RDF namespace and any namespaces used by the enclosed metadata on the `<rdf:RDF>` element, even if these namespaces are already defined on the ancestor elements of the `<rdf:RDF>` element. This increases the re-usability of the RDF block. Furthermore, RDF processing software that does not recognise the CellML namespace can still parse a CellML document, extract the RDF blocks, and perhaps provide useful functionality with the information described in the RDF.

An `<rdf:RDF>` element typically contains one or more `<rdf:Description>` elements, each of which defines an `about` attribute. The value of the `about` attribute must be a valid Uniform Resource Identifier (URI). Metadata may be associated with the *document* it is defined in by assigning the `about` attribute an empty value (`" "`). Metadata may be associated with an element in the current document by defining an attribute of type ID on that element and assigning the `about` attribute on the `<rdf:Description>` element a value equal to the value of that attribute preceded by a hash (`#`). An attribute must be given a type of ID in the document type declaration (DTD) or schema associated with an XML document, and its value must be unique across all

attributes of type ID in a given document. The correct way to do this in a DTD is described in Section 3.3.1 of the XML 1.0 Recommendation.

As was discussed in Section 2.2.1, the `name` attribute that occurs on many CellML elements is not of type ID because it is not necessary that CellML identifiers be unique across a document. To facilitate the association of metadata with CellML elements, a `cmeta:id` attribute in the CellML Metadata namespace may be added to any CellML element. The CellML 1.0 DTD (given in Appendix A.6) declares this attribute to be of type ID for all CellML elements. This declaration prevents CellML elements from having any other attributes of type ID, including attributes in extension namespaces. The MathML 2.0 DTD defines an attribute `id` of type ID for all MathML elements. Extension elements may use their own attributes of type ID, or make use of `cmeta:id` attributes, which CellML processing software is required to treat as if it had type ID. The value of an attribute of type ID must conform to the requirements specified in the XML specification.

For interoperability, an RDF block should be stored in the element about which it contains metadata. This makes the element more re-useable. Elements in the MathML namespace are an exception to this recommendation. The MathML content of a `<cellml:component>` element might be extracted for use in a general MathML processor, which might not be able to handle RDF content. Therefore, metadata on MathML elements should be placed in the containing `<cellml:component>` element. If the RDF block contains metadata about the CellML document, it should be included in the root element of the document. Note that simply putting an RDF block inside an element is not sufficient to indicate that the metadata in the block refers to that element. The `about` attribute on the `<rdf:Description>` element must be used to indicate the resource about which the RDF block contains metadata.

8.3 Examples

Figure 16 demonstrates the use of metadata in CellML. Three RDF blocks are shown: one that provides metadata about the CellML document, one that provides metadata about the model, and one that provides metadata about a component contained in the model. Only the RDF framework elements are shown. The actual metadata is not shown here. Examples in the CellML Metadata Specification will demonstrate how to use the recommended metadata elements.

```
<model
  name="example_metadata_model"
  cmeta:id="model01"
  xmlns="http://www.cellml.org/cellml/1.0#"
  xmlns:cellml="http://www.cellml.org/cellml/1.0#"
  xmlns:cmeta="http://www.cellml.org/metadata/1.0#">

  <!-- This metadata block is about the CellML document -->
  <rdf:RDF
    xmlns:cmeta="http://www.cellml.org/metadata/1.0#"

```

```

xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#">
<rdf:Description about="">
<!-- Some metadata content, such as a last-modified date -->
</rdf:Description>
</rdf:RDF>

<!-- This metadata block is about the CellML model -->
<rdf:RDF
xmlns:cmeta="http://www.cellml.org/metadata/1.0#"
xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#">
<rdf:Description about="#model01">
<!--
Some metadata content, such as a species for which
the model is relevant
-->
</rdf:Description>
</rdf:RDF>

<component name="membrane" cmeta:id="comp01">

<!-- This metadata block is about the membrane component -->
<rdf:RDF
xmlns:cmeta="http://www.cellml.org/metadata/1.0#"
xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#">
<rdf:Description about="#comp01">
<!--
Some metadata content, such as an annotation describing
limitations of this representation of the membrane
-->
</rdf:Description>
</rdf:RDF>

</component>

</model>

```

Figure 16 An example demonstrating how metadata can be embedded in a CellML document using the Resource Description Framework (RDF).

The first RDF block provides metadata about the CellML document. This is indicated by the empty value of the

about attribute on the `<rdf:Description>` element. The second RDF block has a value of `"#model01"` in the about attribute on the `<rdf:Description>` element. This indicates that this metadata provides information about the model that is delimited by the `<cellml:model>` element with a `cmeta:id` attribute value of `"model01"`. The final RDF block provides metadata about the membrane component. This is indicated by the about attribute with a value of `"#comp01"` on the `<rdf:Description>` element.

All three RDF blocks declare the RDF and CellML Metadata namespaces. This makes the RDF blocks portable: the information needed to interpret the RDF will be preserved even if the blocks are extracted from the CellML document.

8.4 Rules for CellML Documents

8.4.1 Proper use of the `cmeta:id` attribute

- A `cmeta:id` attribute (where the `cmeta` prefix is mapped to the CellML Metadata namespace URI defined in Section 2.2.2) may be defined on any CellML element. A `cmeta:id` attribute may also be defined on extension elements for which no attribute of type ID is declared in the DTD, schema or language specification. [On MathML elements, the `mathml:id` attribute must be used. A `cmeta:id` attribute must specifically not be added to MathML elements, because a given element may only contain one attribute of type ID.]

8.4.2 The `<rdf:RDF>` element

1. Allowed use of the `<rdf:RDF>` element

- Any CellML element may contain any number of `<rdf:RDF>` elements. [Metadata may appear on any CellML element, and may be split across multiple `<rdf:RDF>` elements. The recommended practice is to enclose all metadata relevant to a particular resource in a single `<rdf:RDF>` element. In this and subsequent rules, the use of the `rdf` prefix indicates that elements and attributes are in the RDF namespace.]
- The content of an `<rdf:RDF>` element must conform to the Resource Description Framework (RDF) Model and Syntax Specification recommendation from the W3C. [For interoperability, the abbreviated syntax defined in the RDF recommendation should be avoided. However a `rdf:parseType` attribute with a value of `"Resource"` can be added to non-RDF elements to create anonymous resources within an `<rdf:RDF>` element.]

8.5 Rules for Processor Behaviour

8.5.1 Treatment of `cmeta:id` attributes

CellML processing software must treat any `cmeta:id` attributes in a CellML document (where the `cmeta` prefix is mapped to the CellML Metadata namespace URI defined in Section 2.2.2) as if they're of type ID. This has the following consequences for CellML documents:

- A `cmeta:id` attribute must not be defined on a non-CellML element for which the DTD, schema or language specification has already declared an attribute of type ID.
- The values of all `cmeta:id` attributes and any other attributes of type ID in a given CellML document must be unique.
- The values of all `cmeta:id` attributes in a CellML document are potential targets for the values of `about` attributes on `<rdf:Description>` elements.

8.5.2 General meaning of metadata

Metadata may refer to the CellML document, the CellML model, or a specific element within the CellML model. The following list documents the intended meaning of metadata on each of these resources. More detailed information can be found in the CellML Metadata Specification.

- Metadata that refers to the CellML document provides information relevant to the document as a whole, independent from the use of the document to specify a model. Examples of metadata that might appear on a CellML document are last modified date (date on which the document was last edited) and publisher (person or organization distributing the document).
- Metadata that refers to the CellML model provides information relevant to the model as a whole. For instance, the model author is the person who created the complete model, even if some of the components were taken from a shared database and have different authors.
- Metadata that refers to a specific CellML element provides information about that element only. It does not provide information about elements that are contained in the referenced element.

A Using The CellML 1.0 DTD

A.1 Introduction

This section contains some recommendations on the use and referencing of the CellML 1.0 DTD, the full text of which is given in Appendix A.6. The document rules and processor behaviour described in this section are not required in valid CellML documents and from CellML conformant processing software, respectively.

A.2 The CellML DOCTYPE declaration

CellML documents may reference the version of the CellML 1.0 DTD maintained on the CellML website with a DOCTYPE declarations of the following form:

```
<!DOCTYPE model SYSTEM "http://www.cellml.org/cellml/cellml_1_0.dtd">
```

"model" may be changed to match the root element of the CellML document.

CellML document authors may change the value of the system identifier to point to a copy of the DTD cached on a local filesystem if this is available. This specification does not define a public identifier for CellML 1.0. For

interoperability, CellML document authors should not define one unless specifically needed for a parser that makes use of SGML catalog files.

A.3 CellML without a DTD

CellML documents are not required to contain a DOCTYPE declaration. In the absence of a DTD, CellML processing software needs to take into account the following important points:

- All `cmeta:id` attributes are of type ID.
- All `id` attributes on MathML elements are of type ID.

The CellML DTD does not define the default values for CellML attributes that are defined in the specification. With or without a DTD, CellML processing software is responsible for setting these values, if not specified.

A.4 Use of MathML within CellML

The CellML 1.0 DTD can be used in conjunction with the MathML 2.0 DTD to validate any MathML content within a CellML document. The CellML DTD contains a reference to the MathML 2.0 DTD, which is maintained at the World Wide Web Consortium's website, inside a conditional section, which by default is not included in the CellML DTD. CellML document authors may cause the conditional section to be included by setting the value of the `"use_mathml_dtd"` entity to `"INCLUDE"` inside the internal subset of their CellML document, as shown below:

```
<!DOCTYPE model SYSTEM "http://www.cellml.org/cellml/cellml_1_0.dtd" [  
  <!ENTITY % use_mathml_dtd "INCLUDE">  
]>
```

If copies of both the CellML 1.0 DTD and the MathML 2.0 DTD are available on the local filesystem, it is possible to conveniently override the reference to the MathML 2.0 DTD by setting the value of the `"mathml_dtd_path"` entity to the appropriate relative path from the CellML DTD to the MathML DTD, as shown below:

```
<!DOCTYPE model SYSTEM "/my/local/copy/of/the/cellml_1_0.dtd" [  
  <!ENTITY % mathml_dtd_path "'relative/path/to/mathml2.dtd'">  
  <!ENTITY % use_mathml_dtd "INCLUDE">  
]>
```

The MathML 2.0 DTD is not on the CellML website, so DOCTYPE declarations of this form will fail when referencing the DTD on the CellML website. The MathML 2.0 DTDs are available as a ZIP archive, a link to

which is given in Appendix A.6 of the MathML 2.0 Recommendation.

The CellML specification requires a `cellml:units` attribute to be defined on all `<mathml:cn>` elements (where the `cellml` and `mathml` prefixes are mapped to the URIs for the CellML and MathML namespaces, respectively). For this reason, the conditional section of the CellML DTD in which the MathML DTD is referenced also redeclares the list of attributes on the `<mathml:cn>` element to include the `cellml:units` attribute. This will prevent XML parsers from finding errors when validating valid CellML documents against the MathML DTD.

A.5 Treatment of namespaces

In Section 2.2.2, it was suggested that the root element of a CellML document set the default namespace and map the `cellml` prefix to the CellML namespace URI. Use of unprefixed element names (i.e., default namespaces) is encouraged to ensure DTD-based validation will work correctly. The CellML 1.0 DTD defines three optional attributes on every CellML element that can be used to set namespaces where necessary. These are `xmlns`, `xmlns:cellml` and `xmlns:cmeta`.

A.6 The CellML 1.0 DTD

The CellML 1.0 DTD is available at the following URI:

http://www.cellml.org/cellml/cellml_1_0.dtd

The CellML 1.0 DTD does not attempt to declare or reference declarations for the metadata framework elements in the RDF namespace. It also doesn't declare the default values of any of the attributes on the CellML elements.

For convenient reference, the text of the DTD is given below.

```
<!--
```

```
FILE : cellml_1_0.dtd
```

```
CREATED : 22 January 2001
```

```
LAST MODIFIED : 10 August 2001
```

```
AUTHOR : Warren Hedley (w.hedley@auckland.ac.nz)  
The Bioengineering Research Group  
The University of Auckland
```

```
DESCRIPTION : This document contains a DTD corresponding to the syntax rules  
defined in the 10 August 2001 Specification for CellML 1.0. This  
specification is available at
```

<http://www.cellml.org/public/specification/20010810/index.html>

SYSTEM IDENTIFIER : http://www.cellml.org/cellml/cellml_1_0.dtd

COPYRIGHT : (2001) Bioengineering Research Group, University of Auckland.

-->

```
<!ENTITY % use_mathml_dtd "IGNORE">
<![%use_mathml_dtd;[
<!ENTITY % mathml-charent.module "IGNORE">
<!ENTITY % mathml_dtd_path
  "'http://www.w3.org/TR/MathML2/dtd/mathml2.dtd'">
<!ENTITY % mathml_dtd PUBLIC "-//W3C//DTD MathML 2.0//EN"
  %mathml_dtd_path;>
```

```
%mathml_dtd;
```

```
<!ATTLIST %cn.qname;
  %MATHML.Common.attrib;
  %att-type;
  %att-base;
  %att-definition;
  %att-encoding;
  cellml:units CDATA #REQUIRED
>
]]>
```

```
<!ENTITY % cellml_common_attributes "
  xmlns CDATA #IMPLIED
  xmlns:cellml CDATA #IMPLIED
  xmlns:cmeta CDATA #IMPLIED
  cmeta:id ID #IMPLIED
">
```

```
<!ELEMENT model (units | component | group | connection)*>
<!ATTLIST model
  %cellml_common_attributes;
  name CDATA #REQUIRED
>
```

```
<!ELEMENT component (units | variable | reaction | math)*>
<!ATTLIST component
  %cellml_common_attributes;
  name CDATA #REQUIRED
>
```

```
<!ELEMENT variable EMPTY>
<!ATTLIST variable
  %cellml_common_attributes;
  name CDATA #REQUIRED
  public_interface (in|out|none) #IMPLIED
  private_interface (in|out|none) #IMPLIED
  units CDATA #REQUIRED
  initial_value CDATA #IMPLIED
>
```

```
<!ELEMENT connection (map_components | map_variables)+>
<!ATTLIST connection
  %cellml_common_attributes;
>
```

```
<!ELEMENT map_components EMPTY>
<!ATTLIST map_components
  %cellml_common_attributes;
  component_1 CDATA #REQUIRED
  component_2 CDATA #REQUIRED
>
```

```
<!ELEMENT map_variables EMPTY>
<!ATTLIST map_variables
  %cellml_common_attributes;
  variable_1 CDATA #REQUIRED
  variable_2 CDATA #REQUIRED
>
```

```
<!ELEMENT units (unit*)>
<!ATTLIST units
  %cellml_common_attributes;
  name CDATA #REQUIRED
  base_units (yes|no) #IMPLIED
```


>

<!ELEMENT unit EMPTY>

<!ATTLIST unit

 %cellml_common_attributes;

 multiplier CDATA #IMPLIED

 prefix CDATA #IMPLIED

 units CDATA #REQUIRED

 exponent CDATA #IMPLIED

 offset CDATA #IMPLIED

>

<!ELEMENT group (relationship_ref | component_ref)+>

<!ATTLIST group

 %cellml_common_attributes;

>

<!ELEMENT relationship_ref EMPTY>

<!ATTLIST relationship_ref

 %cellml_common_attributes;

 relationship CDATA #REQUIRED

>

<!ELEMENT component_ref (component_ref*)>

<!ATTLIST component_ref

 %cellml_common_attributes;

 component CDATA #REQUIRED

>

<!ELEMENT reaction (variable_ref+)>

<!ATTLIST reaction

 %cellml_common_attributes;

 reversible (yes|no) #IMPLIED

>

<!ELEMENT variable_ref (role+)>

<!ATTLIST variable_ref

 %cellml_common_attributes;

 variable CDATA #REQUIRED

>

```
<!ELEMENT role (math?)>
<!ATTLIST role
  %cellml_common_attributes;
  role (reactant|product|activator|catalyst|inhibitor|modifier|rate) #REQUIRED
  direction (forward|backward|both) #IMPLIED
  delta_variable CDATA #IMPLIED
  stoichiometry CDATA #IMPLIED
>
```

B Scripting functionality in CellML

B.1 Introduction

MathML can be extended by defining new operators, the behaviour of which is not defined in MathML. These operators could be used to call functions implemented in software (i.e., implicitly) or defined using a scripting language (i.e., explicitly). CellML 1.0 does not require processing software to implement support for scripting functionality. This section contains some recommendations on best practices for adding scripting functionality to CellML.

For the purposes of this discussion, it is assumed that scripting functionality will be implemented via function calls. Software that extends the functionality of MathML in other ways should extrapolate these recommendations as appropriate.

B.2 Availability of scripts

Functions referenced in a CellML document that are defined using non-MathML syntax should be identified by, and accessible via, a URI. A function could, for instance, be accessible via HTTP from a database using a CGI script that takes an SQL expression as an argument.

B.3 Embedding scripts in CellML

Functions should be defined within elements in an application-specific extension namespace, if the functions are embedded within CellML documents. If the embedded script is only referenced by a single component, the script should be defined within the corresponding `<cellml:component>` element. If the embedded script is referenced by more than one component, the script should be defined within the `<cellml:model>` element.

B.4 Preferred scripting language

Implementors considering adding scripting functionality to CellML processing software are encouraged to use ECMAScript. It is anticipated that, if scripting functionality is officially endorsed by a future version of CellML, ECMAScript will be the scripting language that processing software is required to implement. ECMAScript is the recommended language because it is simple, standardised, and open source interpreting libraries are freely

available.

ECMA is an international industry association that develops standards in information and communication for the European union. ECMA took the scripting language that Netscape developed for adding interactive content to its web browser and developed the formal language specification ECMA-262 (ECMAScript Language Specification).

B.5 Referencing scripts from MathML

The MathML 2.0 Recommendation defines a `<mathml:csymbol>` element for referencing mathematical symbols and constructs that are not defined by MathML. This element should be used for referencing functions defined using non-MathML syntax from within blocks of MathML. Processing software should take into account the following recommendations regarding the use of the `<mathml:csymbol>` element for this purpose.

B.5.1 The `<mathml:csymbol>` element

1. Recommended use of the `<mathml:csymbol>` element

- The `<mathml:csymbol>` element should only appear as the first child element within a `<mathml:apply>` element. [A `<mathml:csymbol>` element should only be used as an operator, which is "applied" to some arguments.]
- After leading and trailing whitespace is removed, the content of a `<mathml:csymbol>` element must be a valid CellML identifier as discussed in Section 2.2.1. This identifier must accurately represent the external function referenced. [The content of a `<mathml:csymbol>` element should preferably be a human-readable identifier for the external function. If the function is defined using a common scripting or programming language, then this identifier should be the name of the function.]
- Each `<mathml:csymbol>` element should define a `definitionURL` and an `encoding` attribute.

2. Recommended use of the `mathml:definitionURL` attribute

- The value of the `definitionURL` attribute should be a valid URI that identifies a single resource containing the definition of the external function.

3. Recommended use of the `mathml:encoding` attribute

- The value of the `encoding` attribute should indicate to processing software the format of the externally defined function. [In the version of CellML that describes how external functions are to be defined within the CellML framework, the `encoding` attribute will be moved into the CellML namespace and will be required to take a value from a controlled vocabulary.]

B.6 Effects of scripts

Functions must be side-effect free. That is, a function must not assign values to variables that are not local to that function. In particular, functions must not alter the values of their arguments or global variables.

C Advanced Units Functionality

C.1 Introduction

CellML 1.0 lays the foundations of a flexible and robust system for the association of units with variables and constants in cellular models. However, it does not require software to make use of the units information contained in CellML documents. This section presents algorithms and examples demonstrating some of the advanced features related to units that CellML processing software might choose to offer modellers. For interoperability, CellML processing software that includes these features should achieve the same results as the algorithms described here, although the exact implementation may differ.

C.2 Terminology

C.2.1 Equivalence of units references

Two units references are considered equivalent if they satisfy one of the following criteria:

- They reference the same units definition from the standard dictionary.
- They reference the same units definition in the current `<component>` element.
- They reference the same units definition in the current `<model>` element, where that units definition is not superceded by a units definition with the same name in the current `<component>` element.

C.2.2 Dimensional equivalence of units definitions

Two units definitions have dimensional equivalence if, when each is recursively expanded and simplified until left with nothing but products of SI and user-defined base units:

- the expanded form of each units definition consists of the same set of base units, and
- the exponent on each base unit is identical in each expanded units definition.

Algorithms for the expansion and simplification of units definitions are given in Appendix C.3.4 and Appendix C.3.1, respectively.

C.3 Algorithms

C.3.1 Simplification of units definitions

It is frequently convenient to be able to simplify a units definition, where this units definition is the result of the application of some mathematical operator to terms which have units associated with them. These operators include the `<times>`, `<divide>` and `<diff>` operators.

The simplification of a units definition is an iterative process in which the number of other units definitions referenced is systematically reduced. References to units definitions may be removed in the following cases:

- If two units references are equivalent (as defined in Appendix C.2.1) and have exponents with equal and opposite value, then they may be replaced by a reference to `dimensionless`.
- If two units references are equivalent (as defined in Appendix C.2.1) then they may be replaced with a single units reference to the same units, where the exponent associated with that units reference is the sum of the exponents on the original two units references.
- If a units definition references `dimensionless` one or more times in addition to some other units, any references to `dimensionless` may be removed.
- If a units definition references `dimensionless` one or more times and references no other units, the definition may be replaced with `dimensionless`.

The above rules only allow the removal of units references that are equivalent as defined in Appendix C.2.1. This scheme would not allow references to identical units definitions in two different components to be cancelled and removed, because the references would not satisfy the equivalence criteria.

C.3.2 Units-based restrictions on the use of MathML operators

This section describes restrictions on the units associated with a collection of terms to which each MathML operator in the CellML set (defined in Section 4.2.3) can be applied. For instance, the `<mathml:plus>` operator can only be applied to terms that have dimensionally equivalent units. The restrictions invalidate the application of certain MathML operators to certain collections of terms based on their units, and are used in equation dimension checking, as described in Appendix C.3.6.

The restrictions are given in Table 5. The `<mathml:root>`, `<mathml:diff>` and `<mathml:log>` elements all take qualifiers, in addition to operands.

Operator

`<times>`, `<divide>`, `<abs>`, `<floor>`, `<ceiling>`

`<eq>`, `<neq>`, `<gt>`, `<lt>`, `<geq>`, `<leq>`, `<plus>`,
`<minus>`

`<and>`, `<or>`, `<xor>`, `<not>`

`<exp>`, `<ln>`, `<factorial>`, `<sin>`, `<cos>`, `<tan>`,
`<sec>`, `<csc>`, `<cot>`, `<sinh>`, `<cosh>`, `<tanh>`,
`<sech>`, `<csch>`, `<coth>`, `<arcsin>`, `<arccos>`,
`<arctan>`, `<arccosh>`, `<arccot>`, `<arccoth>`,
`<arccsc>`, `<arccsch>`, `<arcsec>`, `<arcsech>`,

Restrictions

There are no restrictions on the units of operands for these operators.

These operators, if applied to more than one operand, require all of their operands to have either equivalent units references, as defined in Appendix C.2.1, or to reference units that have dimensional equivalence, as defined in Appendix C.2.2.

These operators require their operands to have units of `cellml:boolean`, as defined in Section 5.5.2.

These operators require their operands to have units of `dimensionless`.

<arcsinh>, <arctanh>

This is a binary arithmetic operator. Its first operand may have any units, and its second operand must have units of `dimensionless`.

<power>

This is a qualified unary operator. Its operand may have any units. The value of the <degree> qualifier element, if present, must have units of `dimensionless`.

<root>

This is a qualified unary operator. Its operand must have units of `dimensionless`. The value of the <logbase> qualifier element, if present, must have units of `dimensionless`.

<log>

The operand of this operator and the value of the <bvar> qualifier element, if present, may have any units. The value of a <degree> qualifier element within the <bvar> qualifier element, if present, must have units of `dimensionless`.

<diff>

Table 5 The restrictions on the units associated with operands and qualifiers for each MathML operator. All elements in this table are in the MathML namespace.

C.3.3 Applying operators to units definitions

The section defines the units resulting from the application of MathML operators to a collection of terms, each with known units. This is needed for units definition conversion and equation dimension checking, as described in Appendix C.3.5 and Appendix C.3.6, respectively. The units on each of the terms in the collection must satisfy the restrictions defined in Appendix C.3.2 for the current operator.

The full set of units calculation rules are described in Table 6.

Operator

<eq>, <neq>, <gt>, <lt>, <geq>, <leq>,
<and>, <or>, <xor>, <not>

<exp>, <ln>, <log>, <factorial>, <sin>,
<cos>, <tan>, <sec>, <csc>, <cot>, <sinh>,
<cosh>, <tanh>, <sech>, <csch>, <coth>,
<arcsin>, <arccos>, <arctan>, <arccosh>,
<arccot>, <arccoth>, <arccsc>,
<arccsch>, <arcsec>, <arcsech>,

Result Units

The result of these operators has units of
`cellml:boolean`.

The result of these operators has units of `dimensionless`.

<code><arcsinh></code> , <code><arctanh></code>	
<code><plus></code> , <code><minus></code> , <code><abs></code> , <code><floor></code> , <code><ceiling></code>	The result of these operators has the same units as the operands.
<code><times></code>	The result of this operator has units that are the product of the units on the operands. This product may be simplified according to the rules outlined in Appendix C.3.1.
<code><divide></code>	The result of this operator has units that are the quotient of the units on the first and second operands. This quotient may be simplified according to the rules outlined in Appendix C.3.1.
<code><power></code>	The result of this operator has units that are the units on the first operand raised to the power of the second operand. If the first operand has units of <code>dimensionless</code> , the result also has units of <code>dimensionless</code> .
<code><root></code>	The result of this operator has units that are the units on the first operand raised to one over the value of the <code><degree></code> qualifier element (the default value of which is 2.0). If the first operand has units of <code>dimensionless</code> , the result also has units of <code>dimensionless</code> .
<code><diff></code>	The result of this operator has units that are the quotient of the units of the operand over the units of the term in the <code><bvar></code> qualifier element raised to the value of the <code><degree></code> qualifier element inside the <code><bvar></code> qualifier element (the default value of which is 1.0). This quotient may be simplified according to the rules outlined in Appendix C.3.1.

Table 6 The units of the result after applying each MathML operator to a collection of terms, where the units on those terms satisfy the restrictions in Appendix C.3.2. All elements in this table are in the MathML namespace.

C.3.4 Expansion of units definitions

This section presents the recommended algorithm for expanding a given units definition into an expression that relates the defined units to only SI and user-defined base units. This algorithm may be used in the conversion of units definitions and equation dimension checking, algorithms for which are defined in Appendix C.3.5 and Appendix C.3.6, respectively. Examples of the expansion of units definitions are given in Appendix C.4.2.

The specific steps in the algorithm depend on whether the units definition to be expanded is simple or complex,

as defined in Section 5.2.2. Both derivations use recursive methods. At each step, any units that are not base units are replaced with expansions based on the appropriate definition.

Expansion of simple units definitions

The formula relating a variable x_U with units of U (where U are simple units), to a variable x_I with units of u_1 (the subunits referenced by the `units` attribute on the `<unit>` element inside the units definition for U), is given in Equation 9. This is based on the simple units definition formula given in Equation 3.

$$x_U[U] = (m_1 p_1) \left[\frac{U}{u_1} \right] x_I[u_1] + o_1[U] \quad (9)$$

m_1 , p_1 and o_1 correspond to the `multiplier`, `prefix` and `offset` attributes on the `<unit>` element, respectively.

The formula relating a variable with units of u_1 to its subunits u_2 , as referenced in the units definition for u_1 , is given in Equation 10.

$$x_I[u_1] = (m_2 p_2) \left[\frac{u_1}{u_2} \right] x_2[u_2] + o_2[u_1] \quad (10)$$

Equation 10 can be substituted into Equation 9 to give Equation 11.

$$x_U[U] = (m_1 p_1) \left[\frac{U}{u_1} \right] \left((m_2 p_2) \left[\frac{u_1}{u_2} \right] x_2[u_2] + o_2[u_1] \right) + o_1[U] \quad (11)$$

The expression defining each new set of units in terms of its subunits can be substituted into the current expression recursively until an expression is reached that relates x_U to x_n with units u_n , which are SI or user-defined base units. At this point expansion stops, and the resulting expression can be simplified. This simplification combines multiplier, prefix and offset terms and combines their units based on the rules defined in Appendix C.3.3. The result is an expression of the form given in Equation 12.

$$x_U[U] = m_t \left[\frac{U}{u_n} \right] x_n[u_n] + o_t[U] \quad (12)$$

The values of m_t and o_t are given by Equation 13 and Equation 14, respectively.

$$m_t = (m_1 p_1) \dots (m_n p_n) \quad (13)$$

$$o_t = m_1 p_1 (o_2 + m_2 p_2 (o_3 + \dots + m_{n-1} p_{n-1} o_n)) \quad (14)$$

Expansion of complex units definitions

The formula relating a variable x_U with complex units U to a variable x_A with units that are the product of the subunits referenced in the units definition for U is given in Equation 15. This is based on the complex units

definition formula given in Equation 4.

$$x_U[U] = (m_{A1} \dots m_{An} p_{A1}^{e_{A1}} \dots p_{An}^{e_{An}}) \left[\frac{U}{u_{A1}^{e_{A1}} \dots u_{An}^{e_{An}}} \right] x_A[u_{A1}^{e_{A1}} \dots u_{An}^{e_{An}}] \quad (15)$$

The m_{Ai} , p_{Ai} , u_{Ai} and e_{Ai} refer to the values of the `multiplier`, `prefix`, `units` and `exponent` attributes on the i -th `<unit>` element inside the units definition for U , respectively. If, at the first step, the c -th set of units referenced is not SI or user-defined base units, the formula relating u_{Ac} to the c -th set of units referenced to its subunits is given in Equation 16.

$$x_U[u_{Ac}] = (m_{B1} \dots m_{Bn} p_{B1}^{e_{B1}} \dots p_{Bn}^{e_{Bn}}) \left[\frac{U}{u_{B1}^{e_{B1}} \dots u_{Bn}^{e_{Bn}}} \right] x_B[u_{B1}^{e_{B1}} \dots u_{Bn}^{e_{Bn}}] \quad (16)$$

Note that if u_{Ac} is a simple units definition, then the right hand side of Equation 16 will take the form of a simple units definition, but *without the constant offset term*.

An expansion for x_A that incorporates the expansion for u_{Ac} is obtained by multiplying both sides of Equation 16 by unitary coefficients with units of the product of all of the units referenced on the right hand side of Equation 15 with the exception of u_{Ac} . Expansion of units definitions involves constructing expansions for the variable on the right hand side. The expansion continues recursively as long as any of the units on the right hand side of Equation 15 are not SI or user-defined base units. The resulting expansion can be simplified by combining multiplier and prefix terms to give an expression with the form given in Equation 17.

$$x_U[U] = m_t \left[\frac{U}{u_{A1}^{e_{A1}} \dots u_{An}^{e_{An}}} \right] x_A[u_{A1}^{e_{A1}} \dots u_{An}^{e_{An}}] \quad (17)$$

u_{Ii} corresponds to the units referenced by the i -th units definition referenced in the I -th definition to be expanded, and m_t is given by Equation 18.

$$m_t = m_{A1} \dots m_{An} p_{A1}^{e_{A1}} \dots p_{An}^{e_{An}} \quad (18)$$

m_{Ii} , p_{Ii} and e_{Ii} correspond to the values of the `multiplier`, `prefix` and `exponent` attributes on the i -th `<unit>` element in the I -th definition to be expanded.

C.3.5 Conversion between units definitions

This section presents an algorithm that specifies a possible method for converting a variable's value from one set of units to another. An example demonstrating the use of this algorithm is given in Appendix C.4.3.

If the two variable declarations that are to be mapped both reference equivalent units definitions as defined in Appendix C.2.1, then there is a one-to-one mapping between the variable's value in both components.

If the two variable declarations that are to be mapped reference different units definitions, then software may choose to calculate a conversion formula as follows. Given a variable x with units U_x , the value of which is to be passed to a variable y with units U_y , the following steps should be followed:

1. The units definitions for U_x and U_y are fully expanded and simplified according to the algorithm presented in Appendix C.3.4. This yields expressions for x and y in terms of x_n and y_n , the units of which are products of only SI and user-defined base units. The expression for x will be of the form given in Equation 19 if U_x is a simple units definition, or of the form given in Equation 20 if U_x is a complex units definition.

$$x[U_x] = m_x \left[\frac{U_x}{u_x} \right] x_n[u_x] + c_x[U_x] \quad (19)$$

$$x[U_x] = m_x \left[\frac{U_x}{u_x} \right] x_n[u_x] \quad (20)$$

In Equation 19, u_x corresponds to the base units referenced in the full expansion of U_x , whereas in Equation 20, u_x corresponds to the product of all of the base units in the full expansion raised to the appropriate exponents.

2. It should be considered an error if the units for x_n (u_x) and y_n (u_y) do not have equivalent dimensions as defined in Appendix C.2.2.
3. The expansion of x is inverted to give an expression for x_n . The inverted forms of Equation 19 and Equation 20 are given in Equation 21 and Equation 22, respectively.

$$x_n[u_x] = \frac{1}{m_x} \left[\frac{u_x}{U_x} \right] (x[U_x] - c_x[U_x]) \quad (21)$$

$$x_n[u_x] = \frac{1}{m_x} \left[\frac{u_x}{U_x} \right] x[U_x] \quad (22)$$

The appropriate expression for x_n can then be substituted for y_n in the expansion of y . This yields an equation for y in terms of x , which can be used to convert variable values.

C.3.6 Equation dimension checking

This section presents an algorithm that can be used to verify that an equation is consistent with respect to the dimensions of the units definitions referenced by all numbers and variables. An example that demonstrates the process of equation dimension checking for an equation defined in MathML and CellML is given in Appendix C.4.4.

This algorithm relies on the restrictions and behaviour of the different operators with respect to units defined in Appendix C.3.2 and Appendix C.3.3. Future versions of the specification may extend this algorithm to handle other operators. The steps in the algorithm are:

1. The equation is split into a tree of terms, in which each parent term is obtained by the application of a single operator to its children. The root of the tree is the entire equation, which is created by applying a

relational operator (typically the equals operator) to its child terms. All other terms in the tree are created by applying arithmetic operators to child terms.

2. The units definitions for the terms at the leaves of the tree (which will be variables, numbers, or MathML constants elements) are expanded into functions of the SI and user-defined base units, using the algorithm presented in Appendix C.3.4.
3. Starting at the leaves of the tree, sets of child terms are recursively removed from the tree and units assigned to the parent terms. The removal of each set of terms follows the following steps:
 1. The child terms are compared against the restrictions described in Appendix C.3.2 for the current operator. It should be considered an error if they do not satisfy these restrictions, in which case the equation has inconsistent dimensions.
 2. Units are assigned to the parent term as defined in Appendix C.3.3 for the current operator.
4. The equation has self-consistent dimensions if no inconsistencies were found during the recursive removal of child terms during the traversal from leaves to root.

C.4 Examples

C.4.1 User-defined units and new base units

In Figure 17, the example units definitions given in Section 5.3 are reproduced. These examples are used in the subsequent advanced examples.

```
<!-- User-defined Base Units -->
<units name="pH" base_units="yes" />

<!-- Simple Units Definitions -->
<units name="inch">
  <unit multiplier="2.54" prefix="centi" units="metre" />
</units>

<units name="fahrenheit">
  <unit multiplier="1.8" units="celsius" offset="32.0" />
</units>

<!-- Complex Units Definitions -->
<units name="celsius_per_centimetre">
  <unit units="celsius" />
```

```

<unit prefix="centi" units="metre" exponent="-1" />
</units>

<units name="fahrenheit_per_inch">
<unit units="fahrenheit" />
<unit units="inch" exponent="-1" />
</units>

<units name="pH_per_celsius">
<unit units="pH" />
<unit units="celsius" exponent="-1" />
</units>

```

Figure 17 Some examples of the use of the `<units>` element demonstrating the definition of simple and complex units.

C.4.2 Expansion of user-defined units

In this section, the expansion of user-defined units according to the algorithm described in Appendix C.3.4 is demonstrated for each of the units definitions given in Figure 17.

The first `<units>` element in Figure 17 defines units named `pH`, and defines a `base_units` attribute with a value of "yes". This indicates that it should be treated by processing software as if it were an SI base unit, and that it cannot be expanded.

The definition of `inch` in Figure 17 is a simple units definition as it references only a single unit with an exponent of one. When the appropriate terms are substituted into Equation 3, the conversion from `metre` to `inch` is given by Equation 23. `metre` is a SI base unit, so no further expansion is necessary.

$$\begin{aligned}
 x_{\text{new}}[\text{inch}] &= (2.54 \cdot 10^{-2}) \left[\frac{\text{inch}}{\text{metre}} \right] x_{\text{old}}[\text{metre}] \\
 &= 0.0254 \left[\frac{\text{inch}}{\text{metre}} \right] x_{\text{old}}[\text{metre}]
 \end{aligned}
 \quad (23)$$

The definition of `fahrenheit` is in terms of `celsius`, which is an SI derived unit. The expansion from `celsius` to `kelvin`, an SI base unit, is obtained from Section 2.1.1.5 of the SI standard, and is given in Equation 24.

$$x_{\text{new}}[\text{celsius}] = 1.0 \left[\frac{\text{celsius}}{\text{kelvin}} \right] x_{\text{old}}[\text{kelvin}] - 273.15[\text{celsius}] \quad (24)$$

The first step in the expansion of the `fahrenheit` definition is given in Equation 25.

$$x_f [\text{fahrenheit}] = 1.8 \left[\frac{\text{fahrenheit}}{\text{celsius}} \right] x_c [\text{celsius}] + 32.0 [\text{fahrenheit}] \quad (25)$$

The x_c term can be replaced with the expansion of the `celsius` definition from Equation 24, as shown in Equation 26.

$$x_f [\text{fahrenheit}] = 1.8 \left[\frac{\text{fahrenheit}}{\text{celsius}} \right] \left(1.0 \left[\frac{\text{celsius}}{\text{kelvin}} \right] x_k [\text{kelvin}] - 273.15 [\text{celsius}] \right) + 32.0 [\text{fahrenheit}] \quad (26)$$

The 1.0 and 273.15 terms can be multiplied by the 1.8. The units on the resulting terms are the products of the units on the operands, as described in Appendix C.3.3, and these can be simplified according to the rules given in Appendix C.3.1. The final expansion of `fahrenheit` is given in Equation 27

$$\begin{aligned} x_f [\text{fahrenheit}] &= 1.8 \left[\frac{\text{fahrenheit}}{\text{kelvin}} \right] x_k [\text{kelvin}] - 491.67 [\text{fahrenheit}] + 32.0 [\text{fahrenheit}] \\ &= 1.8 \left[\frac{\text{fahrenheit}}{\text{kelvin}} \right] x_k [\text{kelvin}] - 459.67 [\text{fahrenheit}] \end{aligned} \quad (27)$$

`celsius_per_centimetre` is the first of the complex units definitions given in Figure 17. The first step in the expansion of this definition is given by Equation 28, which is obtained by substituting the appropriate terms in Equation 4.

$$x_{\text{new}} [\text{celsius_per_centimetre}] = (10^{-2})^{-1} \left[\frac{\text{celsius_per_centimetre}}{\text{celsius metre}^{-1}} \right] x_{\text{old}} [\text{celsius metre}^{-1}] \quad (28)$$

The expansion of the x_{old} term, which has units that are a product, is not obvious. This term must be expanded to continue. `metre` is an SI base unit, so need not be expanded. However, `celsius` is an SI derived unit, the expansion of which is given in Equation 24. Because this expansion is to be substituted into a complex units definition, the offset term is dropped. The next step in the expansion makes use of the modified `celsius` definition in Equation 29 and the identity in Equation 30.

$$v_{\text{new}} [\text{celsius}] = 1.0 \left[\frac{\text{celsius}}{\text{kelvin}} \right] v_{\text{old}} [\text{kelvin}] \quad (29)$$

$$y_{\text{new}} [\text{metre}^{-1}] = 1.0 \left[\frac{\text{metre}^{-1}}{\text{metre}^{-1}} \right] y_{\text{old}} [\text{metre}^{-1}] \quad (30)$$

The result of multiplying Equation 30 and Equation 29 is given in Equation 31.

$$v_{\text{new}} [\text{celsius}] y_{\text{new}} [\text{metre}^{-1}] = 1.0 \left[\frac{\text{celsius}}{\text{kelvin}} \right] v_{\text{old}} [\text{kelvin}] 1.0 \left[\frac{\text{metre}^{-1}}{\text{metre}^{-1}} \right] y_{\text{old}} [\text{metre}^{-1}] \quad (31)$$

The v_{new} and y_{new} variables can be multiplied together to produce a new unknown z_{new} which has units which are the product of the units on v_{new} and y_{new} . Similarly z_{old} is the product of v_{old} and y_{old} . The result is given in Equation 32, where the scale factors have also been multiplied.

$$z_{\text{new}} [\text{celsius metre}^{-1}] = 1.0 \left[\frac{\text{celsius metre}^{-1}}{\text{kelvin metre}^{-1}} \right] z_{\text{old}} [\text{kelvin metre}^{-1}] \quad (32)$$

z_{new} from Equation 32 can be substituted in place of Equation 28, and the result simplified to give the complete expansion of `celsius_per_centimetre` shown in Equation 33.

$$\begin{aligned} x_{\text{new}} [\text{celsius.per.centimetre}] &= (10^{-2})^{-1} \left[\frac{\text{celsius.per.centimetre}}{\text{celsius metre}^{-1}} \right] x_{\text{old}} [\text{celsius metre}^{-1}] \\ &= 100.0 \left[\frac{\text{celsius.per.centimetre}}{\text{celsius metre}^{-1}} \right] 1.0 \left[\frac{\text{celsius metre}^{-1}}{\text{kelvin metre}^{-1}} \right] z_{\text{old}} [\text{kelvin metre}^{-1}] \quad (33) \\ &= 100.0 \left[\frac{\text{celsius.per.centimetre}}{\text{kelvin metre}^{-1}} \right] z_{\text{old}} [\text{kelvin metre}^{-1}] \end{aligned}$$

The definition of `fahrenheit_per_inch` can be handled in the same way. The first step in the expansion is given by Equation 34.

$$x_{\text{new}} [\text{fahrenheit.per.inch}] = 1.0 \left[\frac{\text{fahrenheit.per.inch}}{\text{fahrenheit inch}^{-1}} \right] x_{\text{old}} [\text{fahrenheit inch}^{-1}] \quad (34)$$

The expansion of x_{old} requires the removal of the offset from the expansion of `fahrenheit` from Equation 27 and the inversion of the expansion of `inch` from Equation 23. These are given in Equation 35 and Equation 36, respectively.

$$v_{\text{new}} [\text{fahrenheit}] = 1.8 \left[\frac{\text{fahrenheit}}{\text{kelvin}} \right] v_{\text{old}} [\text{kelvin}] \quad (35)$$

$$y_{\text{new}} [\text{inch}^{-1}] = 39.370 \left[\frac{\text{inch}^{-1}}{\text{metre}^{-1}} \right] y_{\text{old}} [\text{metre}^{-1}] \quad (36)$$

Multiplying Equation 35 and Equation 36 and simplifying yields Equation 37.

$$z_{\text{new}} [\text{fahrenheit inch}^{-1}] = 70.866 \left[\frac{\text{fahrenheit inch}^{-1}}{\text{kelvin metre}^{-1}} \right] z_{\text{old}} [\text{kelvin metre}^{-1}] \quad (37)$$

z_{new} from Equation 37 can be substituted into Equation 34 in place of x_{old} and the result simplified to give the complete expansion of `fahrenheit_per_inch` in Equation 38.

$$\begin{aligned} x_{\text{new}} [\text{fahrenheit.per.inch}] &= 1.0 \left[\frac{\text{fahrenheit.per.inch}}{\text{fahrenheit inch}^{-1}} \right] x_{\text{old}} [\text{fahrenheit inch}^{-1}] \\ &= 1.0 \left[\frac{\text{fahrenheit.per.inch}}{\text{fahrenheit inch}^{-1}} \right] 70.866 \left[\frac{\text{fahrenheit inch}^{-1}}{\text{kelvin metre}^{-1}} \right] z_{\text{old}} [\text{kelvin metre}^{-1}] \quad (38) \\ &= 70.866 \left[\frac{\text{fahrenheit.per.inch}}{\text{kelvin metre}^{-1}} \right] z_{\text{old}} [\text{kelvin metre}^{-1}] \end{aligned}$$

The first step in the expansion of `pH_per_celsius` is given in Equation 39.

$$x_{\text{new}} [\text{pH.per.celsius}] = 1.0 \left[\frac{\text{pH.per.celsius}}{\text{pH celsius}^{-1}} \right] x_{\text{old}} [\text{pH celsius}^{-1}] \quad (39)$$

When user-defined base units are referenced in a simple or complex units definition, they are treated in the same way as SI base units, and not expanded. The final expansion into a combination of user-defined and SI base units is given in Equation 40.

$$x_{\text{new}} [\text{pH_per_celsius}] = 1.0 \left[\frac{\text{pH_per_celsius}}{\text{pH kelvin}^{-1}} \right] x_{\text{old}} [\text{pH kelvin}^{-1}] \quad (40)$$

C.4.3 Conversion between units definitions

In this section, the algorithm defined in Appendix C.3.5 for converting a variable's value from one set of units to another is presented with respect to a practical example. Figure 18 contains part of a CellML model definition, consisting of two components and one connection. The `legacy_imperial` component defines a variable `x` with units of `fahrenheit_per_inch`. The `modern_si` component defines a variable `y` with units of `celsius_per_centimetre`. A connection between the two components maps `x` to `y`.

```
<component name="legacy_imperial">
  <variable name="x" public_interface="out" units="fahrenheit_per_inch" />
</component>

<component name="modern_si">
  <variable name="y" public_interface="in" units="celsius_per_centimetre" />
</component>

<connection>
  <map_components component_1="legacy_imperial" component_2="modern_si" />
  <map_variables variable_1="x" variable_2="y" />
</connection>
```

Figure 18 In this model fragment, a connection maps a variable `x` with units of `fahrenheit_per_inch` to a variable `y` with units of `celsius_per_centimetre`.

The CellML definitions of both `fahrenheit_per_inch` and `celsius_per_centimetre` are given in Figure 17. It was shown how to obtain expressions that relate each of these units definitions to the SI base units in Appendix C.4.2. These expressions are reproduced in Equation 41 and Equation 42, respectively.

$$x_{\text{psi}} [\text{fahrenheit_per_inch}] = 70.866 \left[\frac{\text{fahrenheit_per_inch}}{\text{kelvin metre}^{-1}} \right] w_{\text{kpm}} [\text{kelvin metre}^{-1}] \quad (41)$$

$$y_{\text{cpm}} [\text{celsius_per_centimetre}] = 100.0 \left[\frac{\text{celsius_per_centimetre}}{\text{kelvin metre}^{-1}} \right] w_{\text{kpm}} [\text{kelvin metre}^{-1}] \quad (42)$$

The value of `x` is transferred to the variable `y` in the mapping. Therefore an equation expressing

celsius_per_centimetre in terms of fahrenheit_per_inch is needed. This can be obtained by rearranging Equation 41 for w_{kpm} , substituting the resulting expression in place of z_{kpm} in Equation 42, and simplifying the result according to the rules defined in Appendix C.3.1. This gives the expression in Equation 43.

$$\begin{aligned}
 y_{cpm} [\text{celsius_per_centimetre}] &= 100.0 \left[\frac{\text{celsius_per_centimetre}}{\text{kelvin metre}^{-1}} \right] z_{kpm} [\text{kelvin metre}^{-1}] \\
 &= \frac{100.0 \left[\frac{\text{celsius_per_centimetre}}{\text{kelvin metre}^{-1}} \right]}{70.866 \left[\frac{\text{fahrenheit_per_inch}}{\text{kelvin metre}^{-1}} \right]} x_{fpi} [\text{fahrenheit_per_inch}] \quad (43) \\
 &= 1.411 \left[\frac{\text{celsius_per_centimetre}}{\text{fahrenheit_per_inch}} \right] x_{fpi} [\text{fahrenheit_per_inch}]
 \end{aligned}$$

C.4.4 Equation dimension checking

In this section, the algorithm defined in Appendix C.3.6 for checking that an equation has consistent units is presented with respect to a practical example. Figure 19 contains the definition of a CellML component sodium_channel_m_gate. This component defines three sets of units (per_millisecond, millivolt, and per_millivolt), two variables (V and alpha_m) and an equation that calculates the value of alpha_m.

```

<component name="sodium_channel_m_gate">
  <units name="per_millisecond">
    <unit prefix="milli" units="second" exponent="-1" />
  </units>
  <units name="millivolt">
    <unit prefix="milli" units="volt" />
  </units>
  <units name="per_millivolt">
    <unit prefix="milli" units="volt" exponent="-1" />
  </units>

  <variable name="V" units="millivolt" />
  <variable name="alpha_m" units="per_millisecond" />

  <math xmlns="http://www.w3.org/1998/Math/MathML">
    <apply><eq />
    <ci> alpha_m </ci>
    <apply><times />
    <cn cellml:units="per_millisecond"> 1.0 </cn>
    <apply><divide />
    <apply><times />
    <cn cellml:units="per_millivolt"> 0.1 </cn>
  </math>

```



```

<apply><plus />
<ci> V </ci>
<cn cellml:units="millivolt"> 25.0 </cn>
</apply>
</apply>
<apply><minus />
<apply><exp />
<apply><times />
<cn cellml:units="per_millivolt"> 0.1 </cn>
<apply><plus />
<ci> V </ci>
<cn cellml:units="millivolt"> 25.0 </cn>
</apply>
</apply>
</apply>
<cn cellml:units="dimensionless"> 1.0 </cn>
</apply>
</apply>
</apply>
</apply>
</math>
</component>

```

Figure 19 Parts of the CellML definition of the Hodgkin Huxley squid axon model. The `sodium_channel_m_gate` component defines three sets of units, two variables `V` and `alpha_m`, and the calculation of `alpha_m`.

Equation 44 gives the equation in Figure 19, where units have been omitted. Equation 45 gives the same equation, with the units associated with each number and variable are included. The first 1.0 in the equation is included specifically for units consistency. It would be possible to associate more complex units with the 0.1 in the numerator of the equation, but this would not accurately reflect the intent of the original model authors. CellML processing software is free to find this and similar terms that do not affect the mathematics and ignore them when interpreting the equation.

$$\alpha_{m} = 1.0 \left(\frac{0.1 (V + 25.0)}{\exp(0.1 (V + 25.0)) - 1.0} \right) \quad (44)$$

$$\alpha_{m} [\text{per_millisecond}] = 1.0 [\text{per_millisecond}] \cdot \left(\frac{0.1 [\text{per_millivolt}] (V [\text{millivolt}] + 25.0 [\text{millivolt}])}{\exp(0.1 [\text{per_millivolt}] (V [\text{millivolt}] + 25.0 [\text{millivolt}])) - 1.0 [\text{dimensionless}]} \right) \quad (45)$$

The first step in the algorithm proposed in Appendix C.3.6 for verifying that a given equation has consistent

dimensions is to convert the equation into a tree of equation parts. Equation 45 can be broken up into the tree shown in Figure 20.

the tree form of `alpha_m` equation

Figure 20 The tree form of Equation 45, in which each non-leaf node is obtained by the application of a single operator to its children. Each operator is represented by a number in square brackets. The effect of each operator is discussed in the text.

At each branch in the tree, a single operator is applied to the child nodes, combining them into a larger parent equation part. Each equation part in the tree has units associated with it. In the case of leaf nodes, these units are obtained from the MathML equation definition. Parent nodes have units defined by the operator and the units on their child nodes, as described in Appendix C.3.3. Dimension checking begins at the leaf nodes, which are recursively removed as the units are evaluated for their parent nodes, which in turn become leaf nodes, as described in Appendix C.3.6.

In the equation tree diagram in Figure 20, the application of each operator to a set of child nodes is denoted by a number in square brackets, where the number reflects the order in which the operations are processed. These operations are discussed below.

1. The `<plus>` operator combines the variable `"V"` and the number `"25.0"` that occur in the denominator of the fraction in Equation 45. The `<plus>` operator requires all of its operands to be dimensionally equivalent, as described in Appendix C.2.2. The resulting equation part will have the same units as the operands. Both `"V"` and `"25.0"` have units of `millivolt`, and so the parent equation part `"V + 25.0"` also has units of `millivolt`.
2. The `<times>` operator (which is not explicitly rendered in Equation 45) multiplies the `"V + 25.0"` term that is the result of the first operation by the number `"0.1"` which has units of `per_millivolt`. The `<times>` operator can be applied to any operands, independent of their units, and the resulting equation part has units that are the product of the units on the operands. In this case, the resultant `"0.1 (V + 25.0)"` term has units that are the product of `millivolt` and `per_millivolt`, which simplifies to `dimensionless` (as described in Appendix C.3.1).
3. The `<exp>` operator is a unary arithmetic operator and its operand must have units of `dimensionless`. The result of applying the operator, in this case `"exp (0.1 (V + 25.0))"`, also has units of `dimensionless`.
4. The `<minus>` operator subtracts the number `"1.0"`, which has units of `dimensionless`, from the term resulting from operation 3. The `<minus>` operator requires both its operands to have the same units and the result assumes those units.
5. The `<plus>` operator is applied to the variable `"V"` and the number `"25.0"` from the numerator of the fraction in Equation 45. Units are handled as in operation 1.
6. Exactly as in operation 2, where the operands are now in the numerator of the fraction in Equation 45.

7. The `<divide>` operator is applied to the results of operations 6 and 4, which both have units of `dimensionless`. The `<divide>` operator can be applied to any operands, independent of their units, and the result has the quotient of the units on the operands. In this case, the resulting fraction has units of `dimensionless`.
8. The `<times>` operator is applied to the number "1.0", which has units of `per_millisecond` and the result of operation 7, which has units of `dimensionless`. The resulting term has units of `per_millisecond`.
9. Finally, the `<equals>` operator is applied to the variable "*alpha_m*" and the term resulting from operation 8. The `<equals>` operator requires that its operands have dimensionally equivalent units.

D Changes

D.1 Changes between 18 May 2001 Final Draft and the 10 August 2001 Recommendation

The full list of changes made between the 18 May 2001 Final Draft and the 10 August 2001 Recommendation of the CellML 1.0 specification can be found on the 18 May 2001 CellML 1.0 Specification Errata. Unfortunately, the full list is too large to include here.

D.2 Changes between 2 March 2001 Draft and the 18 May 2001 Final Draft

The complete list of changes made between the 2 March 2001 Draft and the 18 May 2001 Final Draft of the CellML 1.0 specification can be found at Appendix D of the 18 May 2001 Final Draft.

© 2001-2012 - The CellML Project.